

ECMWF Talk – March 2019

Running Serverless HPC Workloads on Top of Kubernetes and Jupyter Notebooks

Christopher Woods

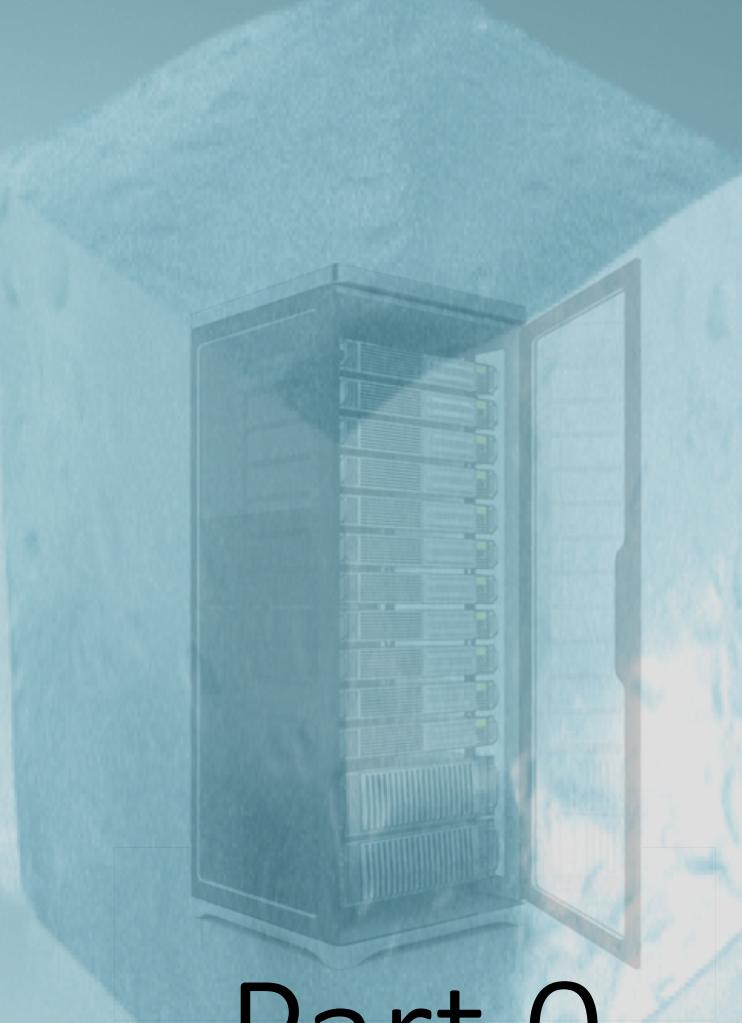
EPSRC Research Software Engineering Fellow

Advanced Computing Research Centre

University of Bristol

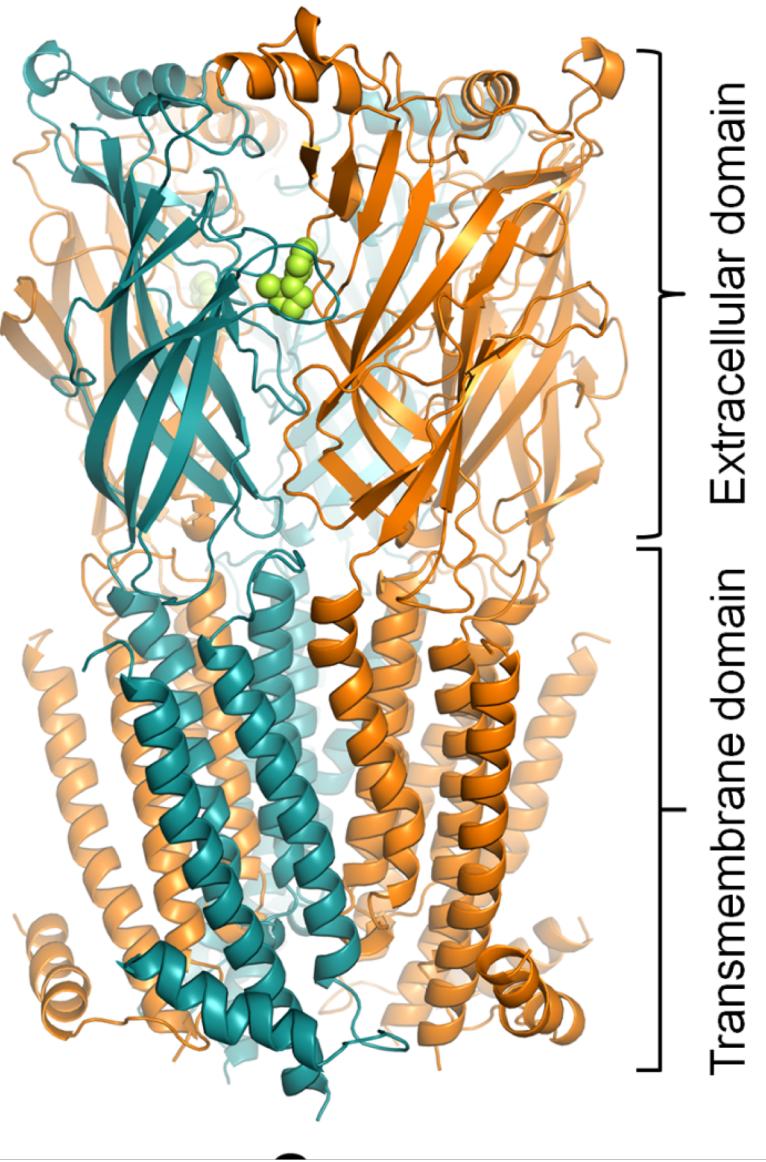


University of
BRISTOL



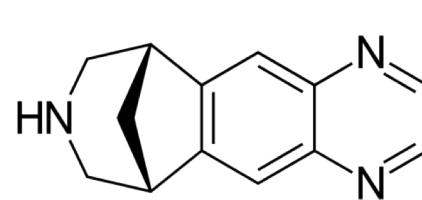
Part 0

Cluster on Premise

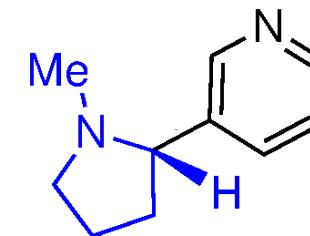


Nicotinic Receptor (Ligand Gated Ion Channel)

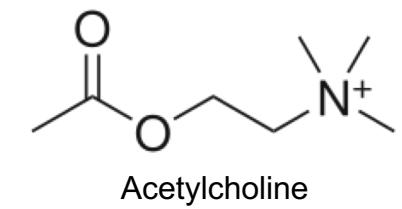
Important drug target for treatment of
Alzheimer's, Schizophrenia, Pain etc.



Varenicline



Nicotine

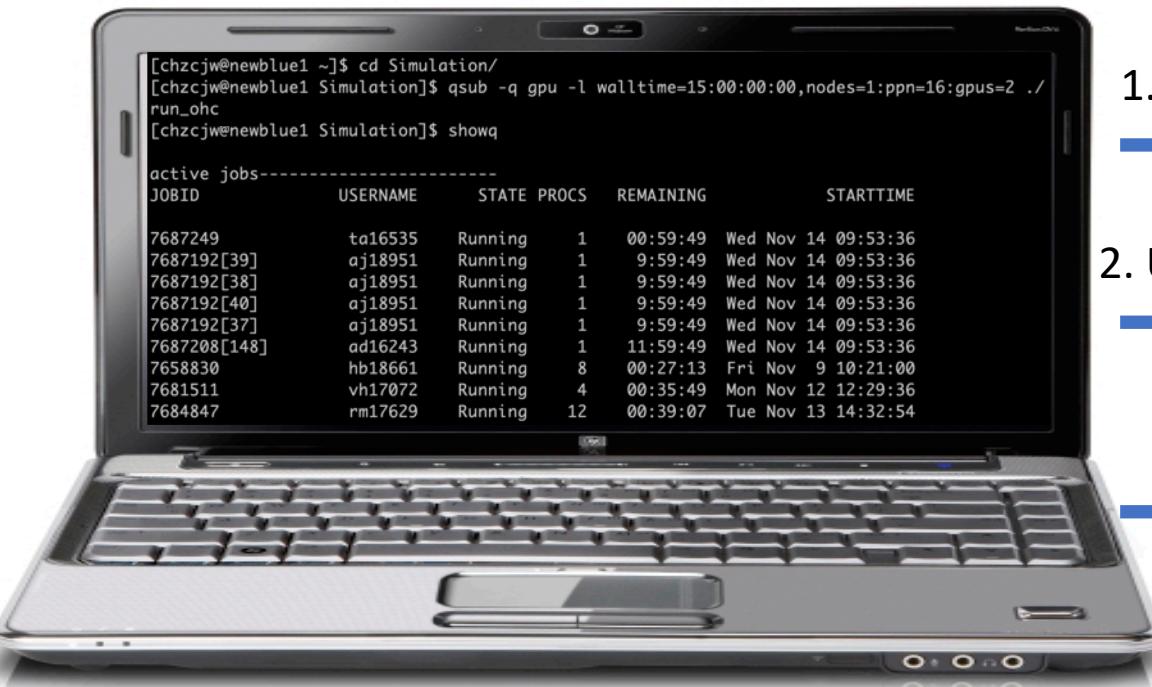


Acetylcholine

To understand this protein, we need to use computer simulations to ask how does the receptor move?

How does the signal move from one side of the receptor to the other?

How would we run these simulation using our on-premise cluster?

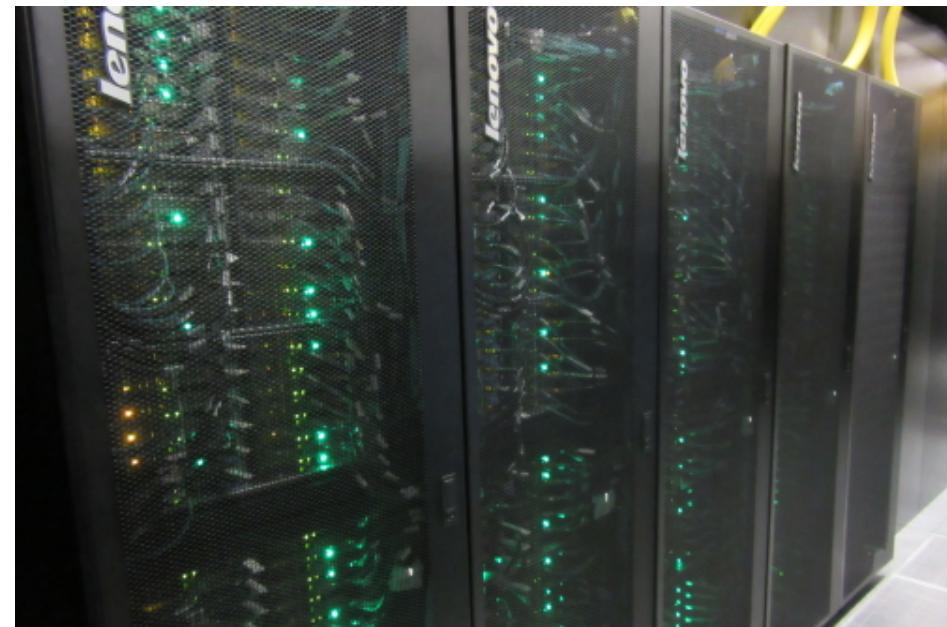


1. SSH to login node

2. Upload input (rsync)

3. Submit job to Q

6. Analyse results



4. Wait in Q

5. Wait for job

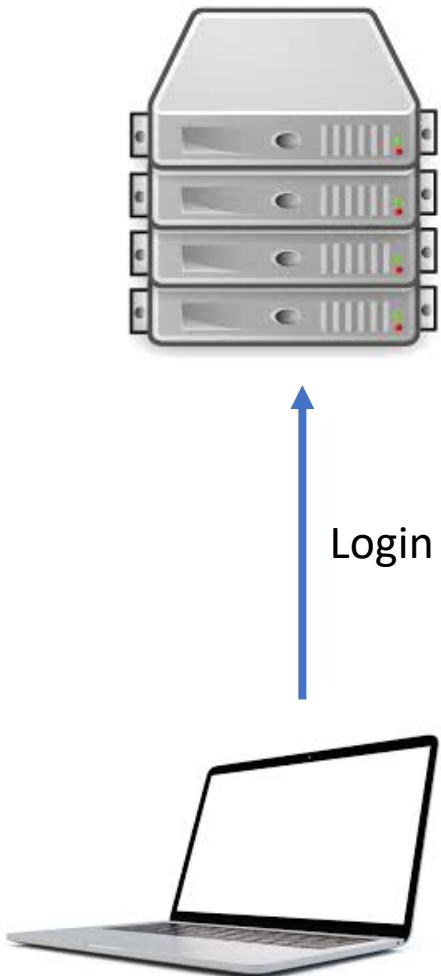
5. Download output (rsync)



Part 1

Cluster in the Cloud

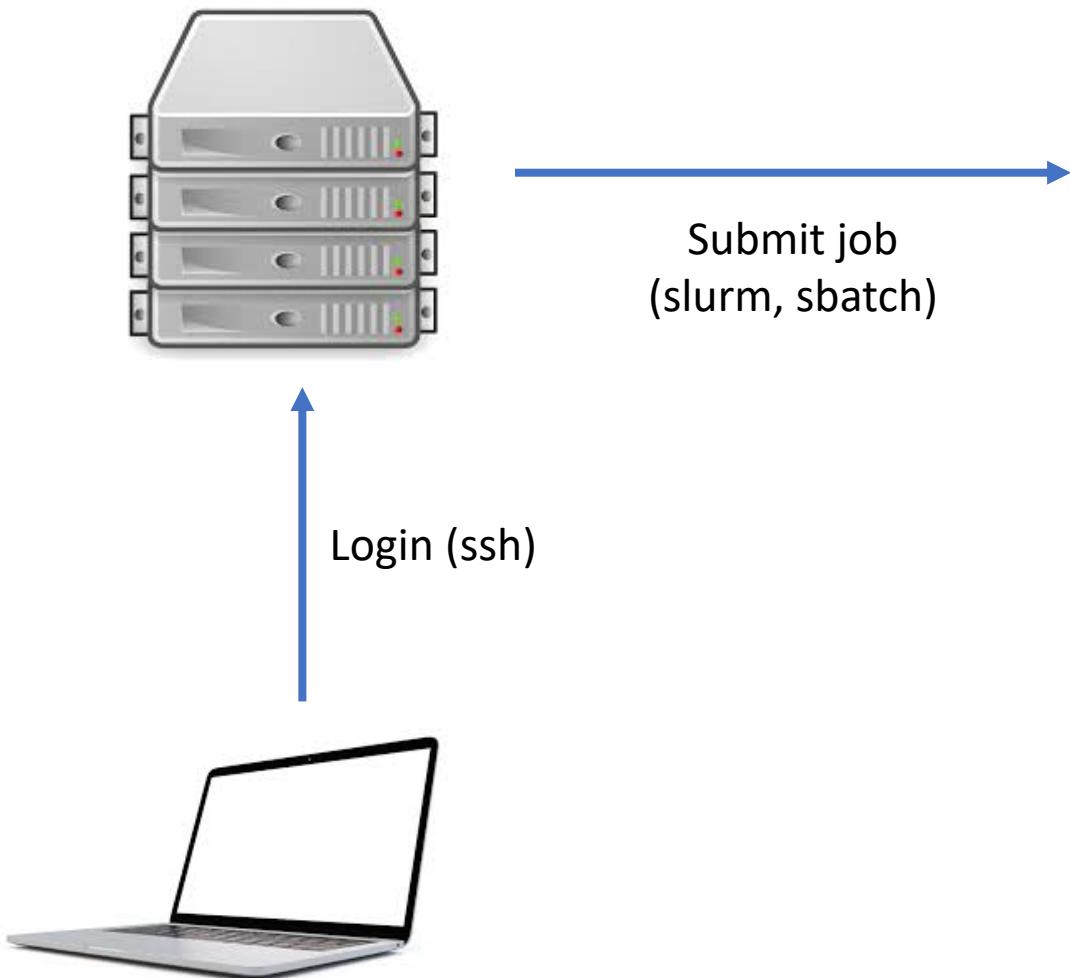
Cloud Login / management VM



<https://cluster-in-the-cloud.readthedocs.io>



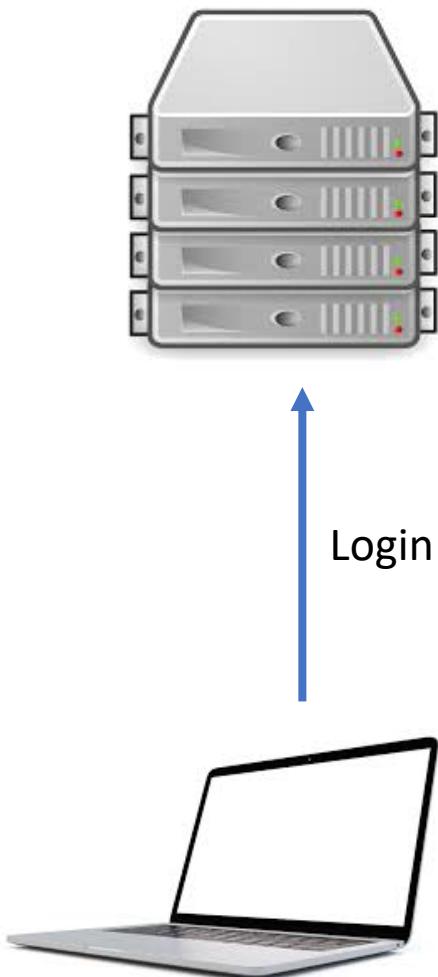
Cloud Login / management VM



<https://cluster-in-the-cloud.readthedocs.io>

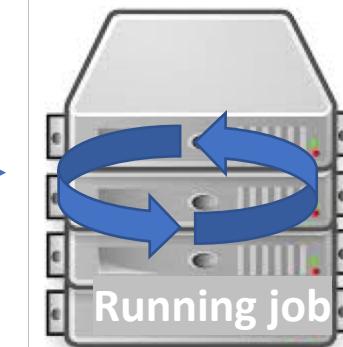


Cloud Login / management VM



Login (ssh)

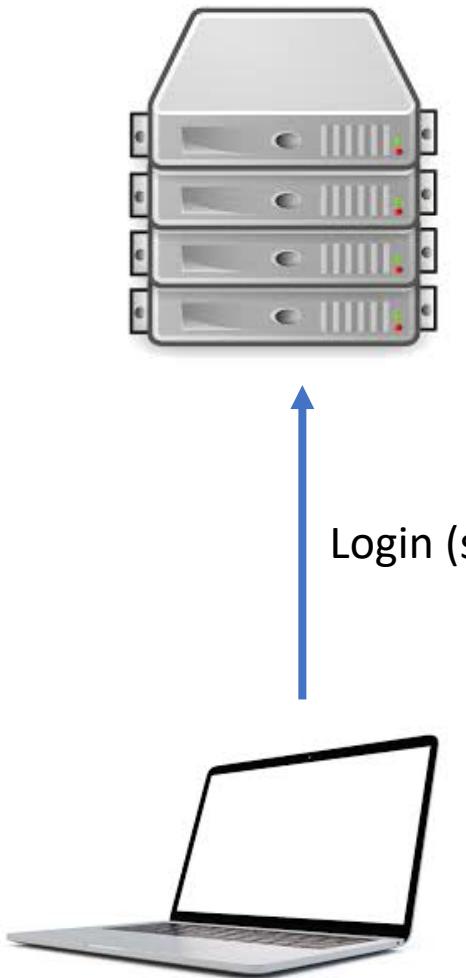
Bare metal Standard2.52



Submit job
(slurm, sbatch)

On

Cloud Login / management VM



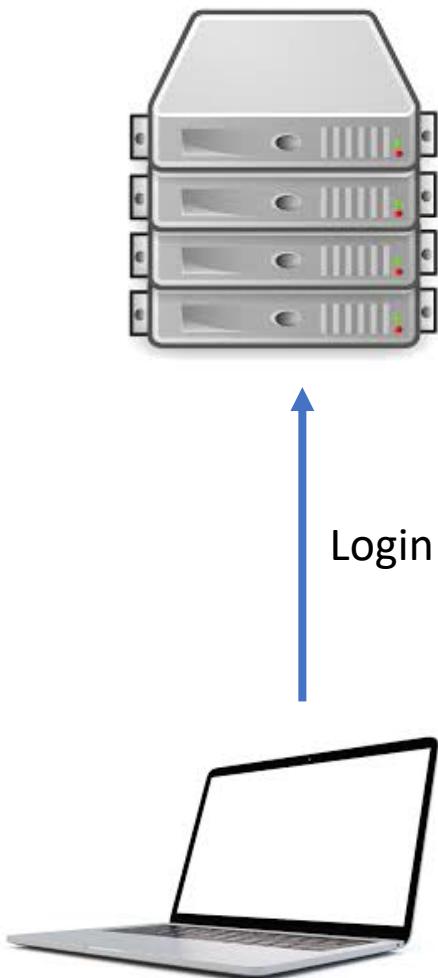
Bare metal Standard2.52



<https://cluster-in-the-cloud.readthedocs.io>

ORACLE®
Cloud Infrastructure

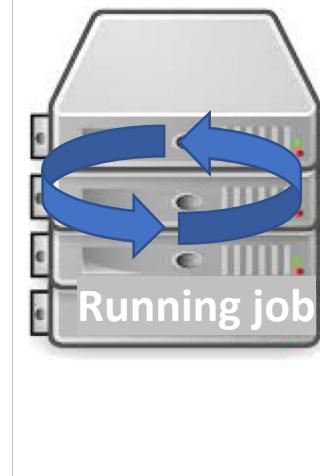
Cloud Login / management VM



ORACLE®
Cloud Infrastructure

Submit job
(slurm, sbatch)

Bare metal Standard2.52

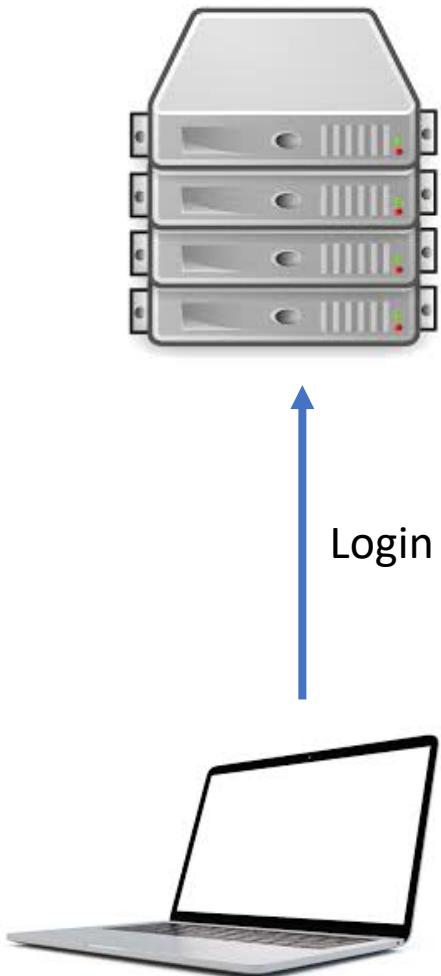


On



On

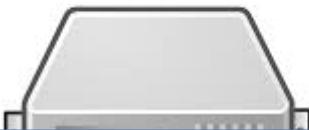
Cloud Login / management VM



<https://cluster-in-the-cloud.readthedocs.io>



Cloud Login / management VM



Cluster in the Cloud

Scales to the size of the job as needed

Only pay for what you use

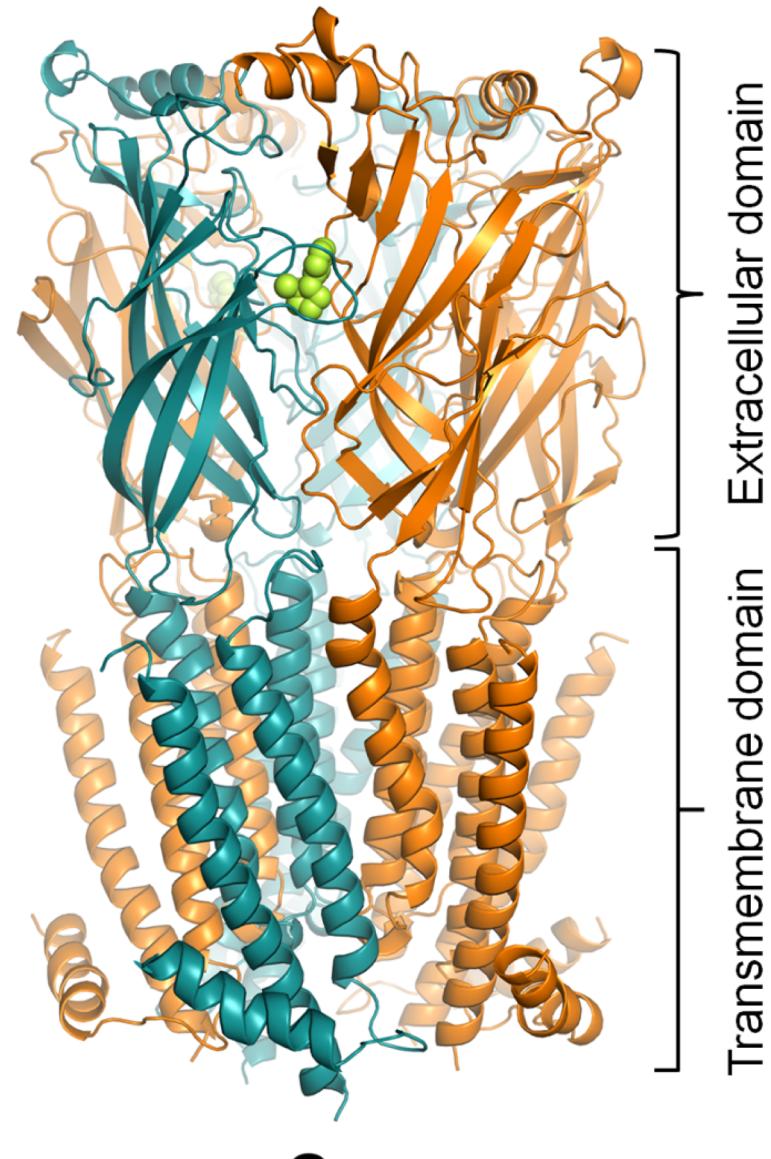
Looks and feels like a “traditional” HPC cluster

Shallow learning curve for researchers



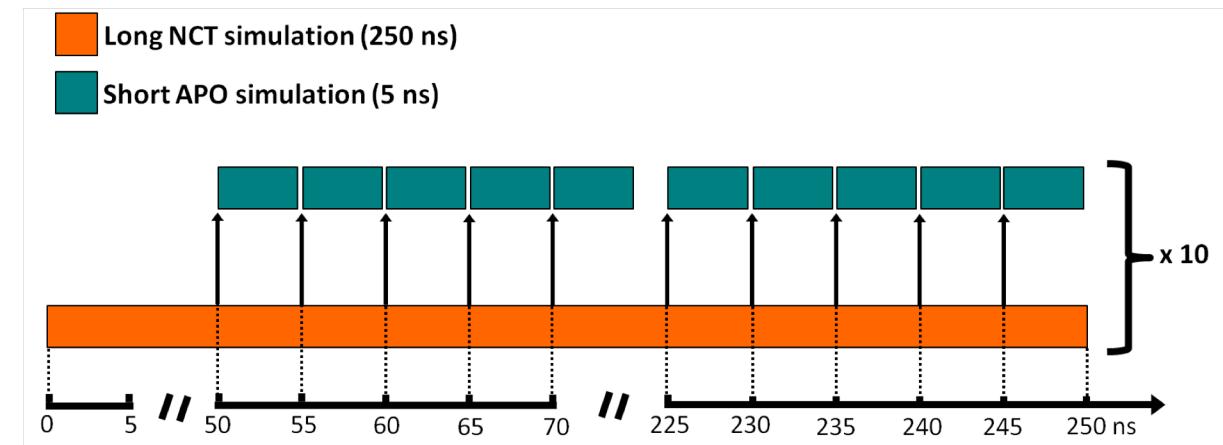
<https://cluster-in-the-cloud.readthedocs.io>





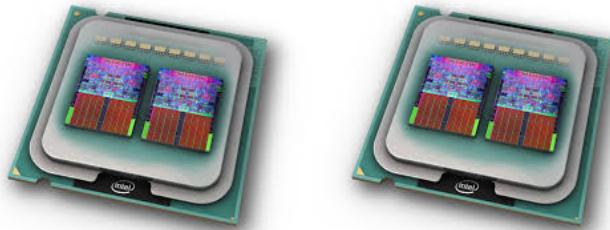
Nicotinic Receptor (Ligand Gated Ion Channel)

Important drug target for treatment of
Alzheimer's, Schizophrenia, Pain etc.



Investigate the signal conduction pathway via 1350 x
5ns molecular dynamics simulations (6.75 μ s total)

On-premise (BlueCrystal 4)



Installed early 2017
2 x Intel Xeon (Broadwell) 14 core

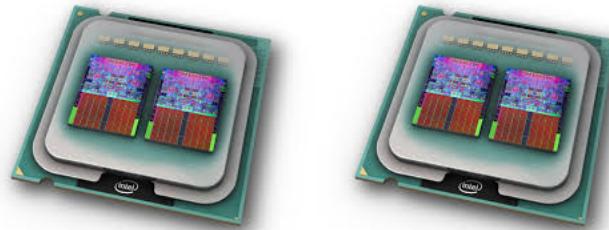
5ns (Gromacs) takes 17 hours

>530 compute nodes shared
between ~1000 users

At most ~10 jobs could run at once

Calculation would take ~90 days

Oracle Cloud (BM.Standard2.52)



Bare Metal Node using latest Platinum Xeons
2 x Intel Xeon (Skylake) 26 core

5ns (Gromacs) takes 8 hours

100 nodes available for 1 user

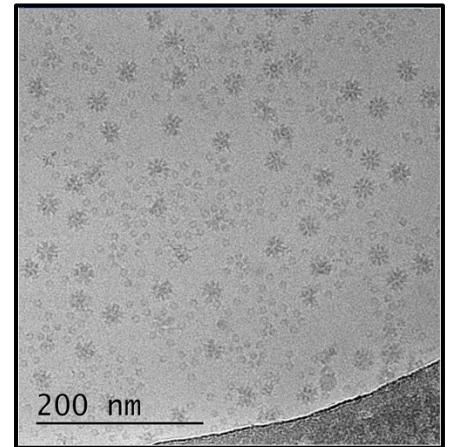
At most 100 jobs could run at once

Calculation took 5 days (Wed-Sun)



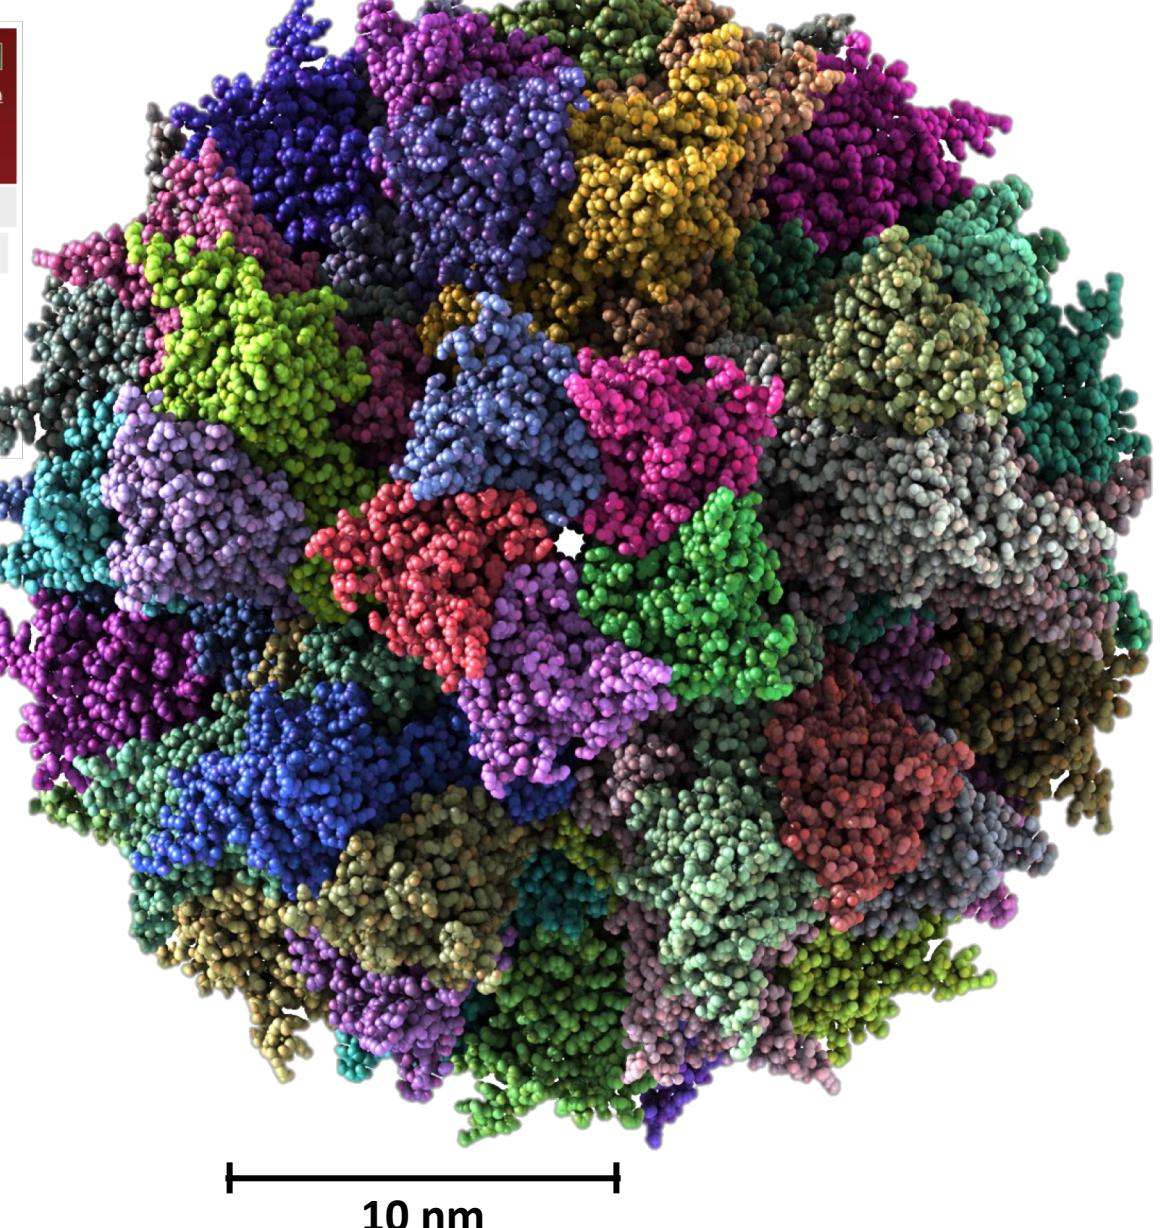
Cryo-electron microscopy wins chemistry Nobel

Jacques Dubochet, Joachim Frank and Richard Henderson share the prize for developing a technique to image biomolecules.



Terabytes of raw data

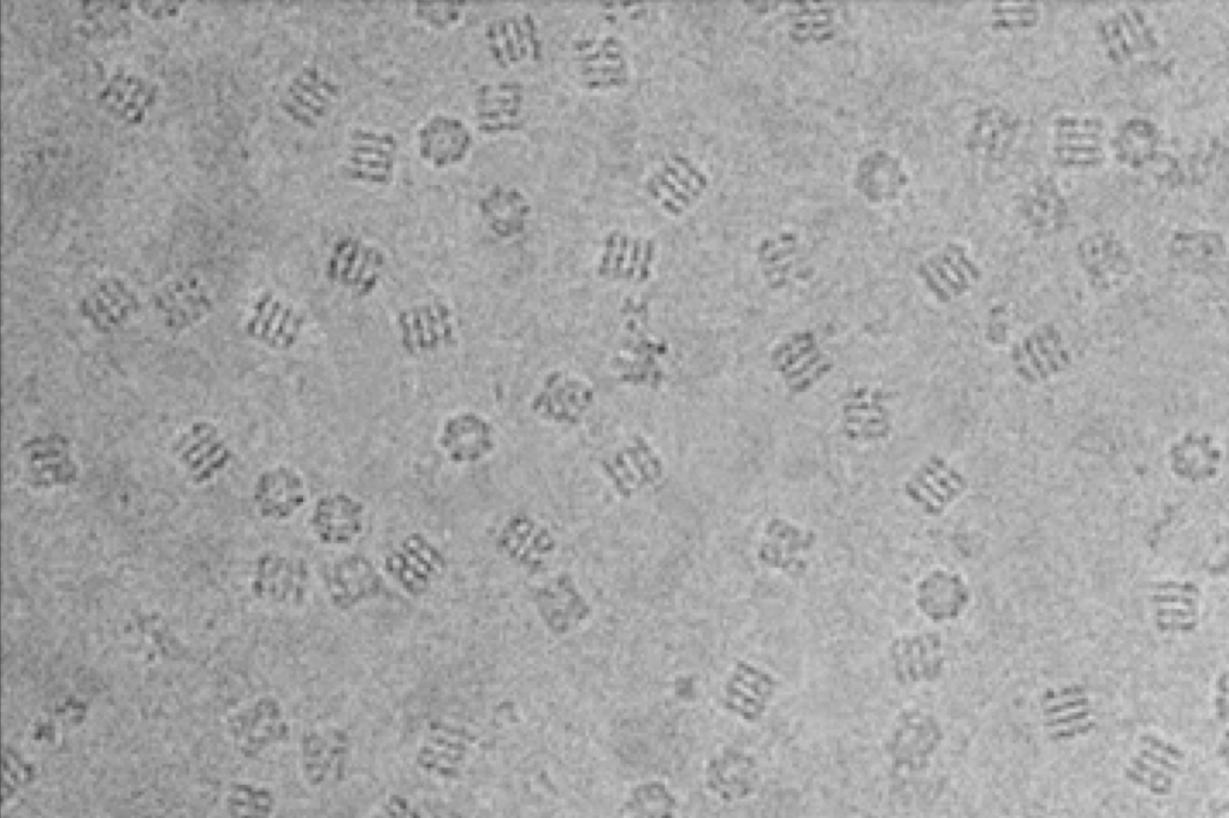
LOTS of
processing



Cryo-electron Microscope

3D atomic structure of the protein

This is an example of an image we collect using the electron microscope



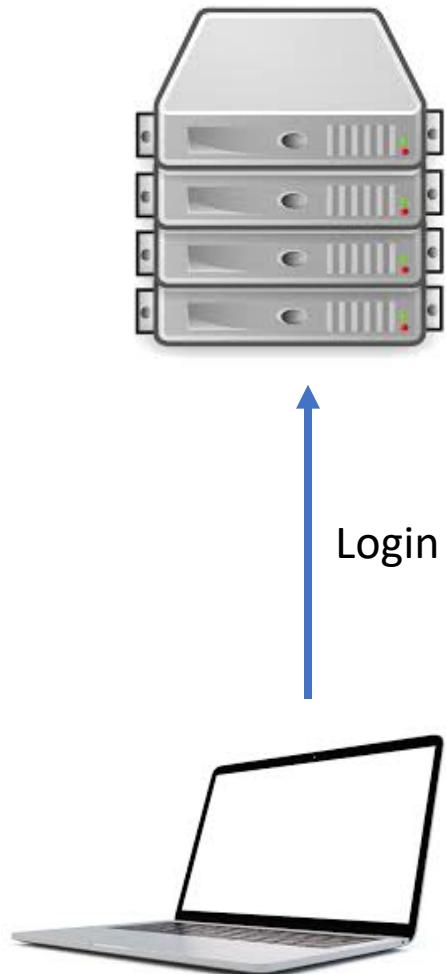
We can see the molecules as they are trapped in the ice, in different orientations

<https://www.youtube.com/watch?v=BJKkC0W-6Qk>

[Gabe Lander](#) This is a very brief introduction to how transmission cryo-electron microscopy is used to solve the structures of protein complexes. Download slides as PDF here: [http://www.ocf.berkeley.edu/~glander/...](http://www.ocf.berkeley.edu/~glander/)

Cryo-EM image analysis (Relion)

Cloud Login / management VM



CPU Node (Stage 1 Relion)

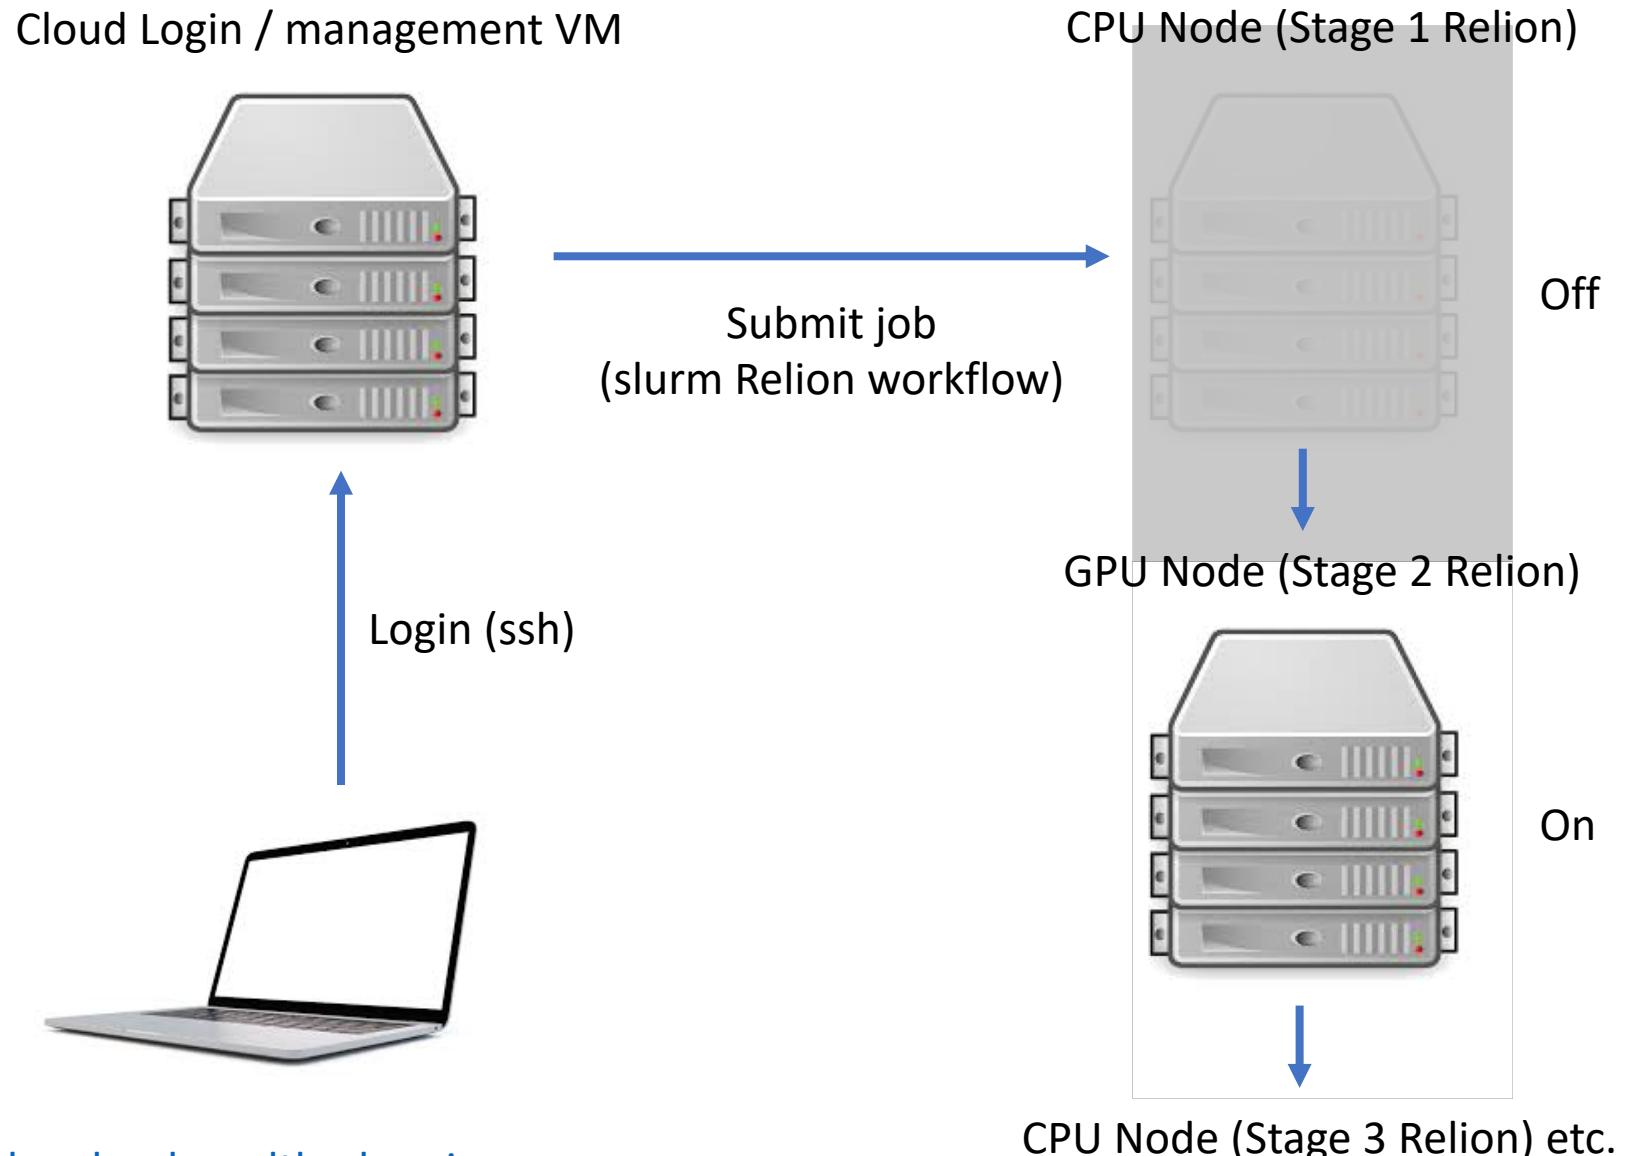


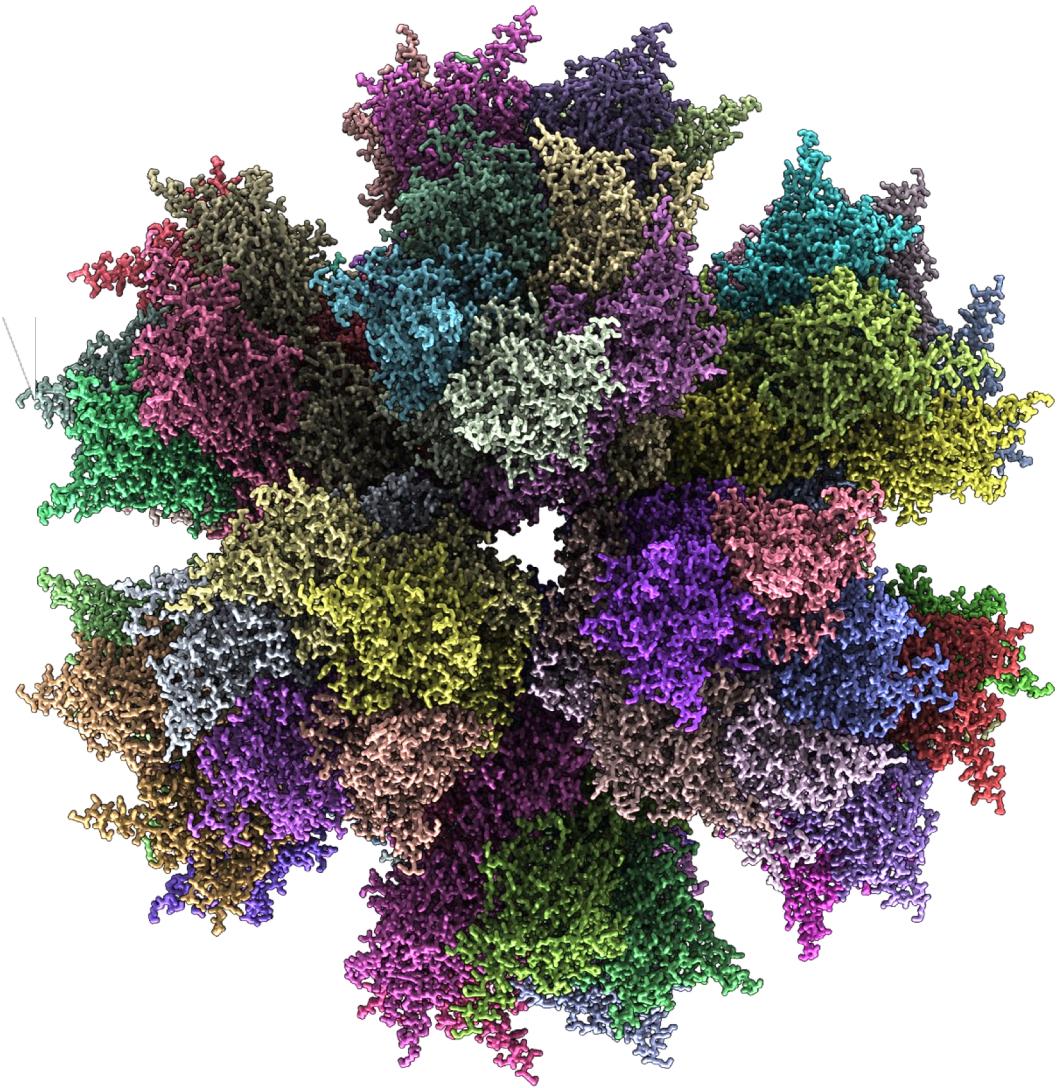
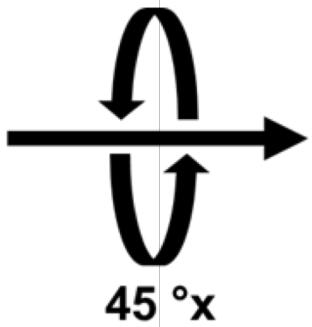
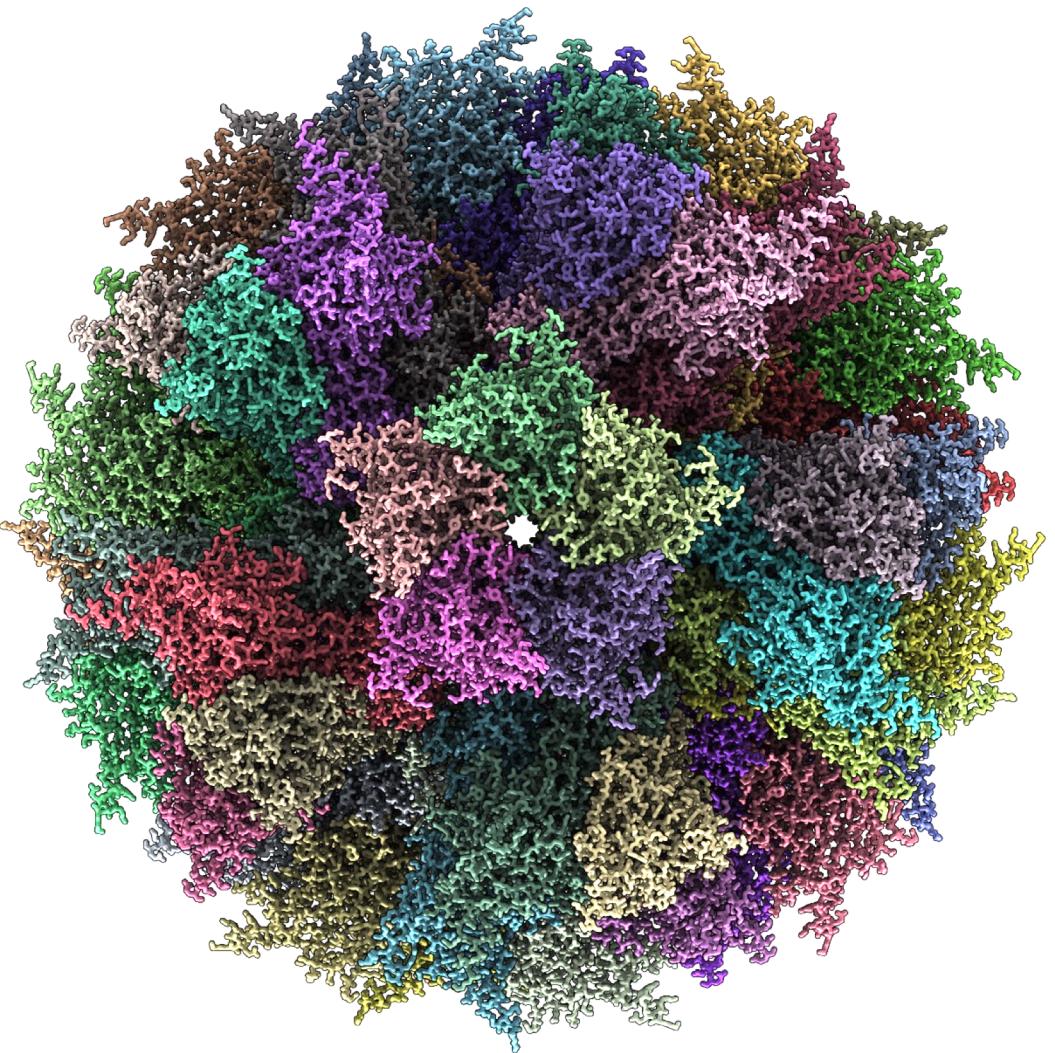
Submit job
(slurm Relion workflow)

On

Login (ssh)

Cryo-EM image analysis (Relion)





Processing of Cryo-EM movie frames to reveal ADDomer atomic structure for structure-based drug design

Understanding the atomic structure enables simulation of vaccine behaviour against new viruses

Will provide information enabling rapid response to new virus mutations

Generating these structures cost just £217 on OCI

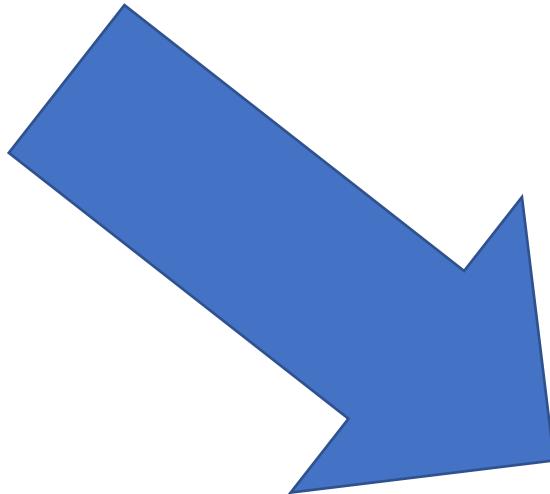
Cluster-in-the-cloud projects

- Completed
 - Molecular dynamics of proteins to investigate cell signaling (5200 cores)
 - Cryo-EM image analysis (TB of data)
 - GPU-accelerated binding affinity prediction (lots of GPUs)
- In progress...
 - Designing silicon devices for detection and control of radiotherapy doses (thousands of quantum physics simulations of x-ray photon trajectories)
 - Detection of dark matter signatures (deep physics)
- ...and many more at other universities ☺

A good waypoint – but not the final destination...



On-premise HPC



Cluster-in-the-cloud

Cloud is more than just a Personal Cluster...

- We gave our researchers access to their own personal “cluster in the cloud”
- This let them perform their experiments in a familiar environment
- But – they were queuing jobs onto their own personal cluster!
- Revolutionary potential of the cloud will be realized if we treat it as an on-demand resource, not just as a “batch-cluster in the sky”



2000



2018

A cellphone is more than just a mobile telephony device

Part 2

Acquire



Interactive molecular dynamics

BioSimSpace is a great tool for playing around with molecular simulations directly and interacting with them in real-time. In this notebook you'll learn how to use BioSimSpace to set up and run an equilibration protocol, then query the running process for information, plot graphs of the latest data, visualise molecular configurations, and analyse trajectory data.

Before we get started, let's import BioSimSpace so that it's available inside of our notebook.

```
In [ ]: import BioSimSpace as BSS
```

Creating a molecular system

First of all we need to load a molecular system.

```
In [ ]: system = BSS.IO.readMolecules(["amber/ala/ala.crd", "amber/ala/ala.top"])
```

We have now created a molecular system. The system consists of an alanine dipeptide molecule in a box of water. To show the number of molecules in the system, run:

```
In [ ]: system.nMolecules()
```

Defining a simulation protocol

BioSimSpace provides functionality for defining various simulation protocols. In this notebook we will construct a typical simulation workflow that uses a sequence of simple protocols, with the output of one forming the input of the next:

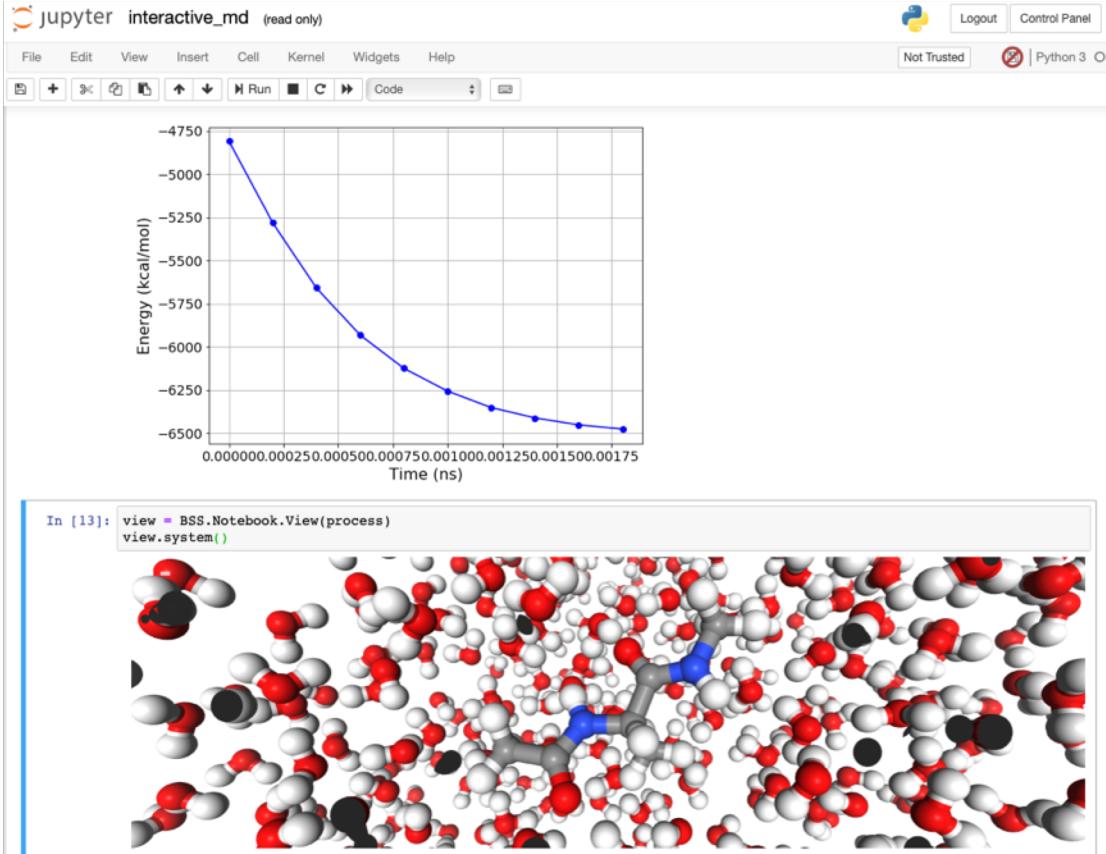
1. *Minimisation*: Energy minimisation the molecular system.
2. *Equilibration*: Equilibration of the system to a target temperature.
3. *Production*: Regular molecular dynamics, run at fixed temperature.
4. *Custom*: A user defined protocol, e.g. a config file for a molecular dynamics package.

When defining a protocol we are configuring the type of simulation that we wish to run, as well as any options for the particular simulation. For example, to create a default equilibration protocol:

```
protocol = BSS.Protocol.Equilibration()
```

This defines a 0.2 nanosecond equilibration protocol at a temperature of 300 Kelvin. For convenience, let's reduce the runtime. We'll also perform a heating protocol and will restrain the position of atoms in the backbone.

Jupyter Notebooks

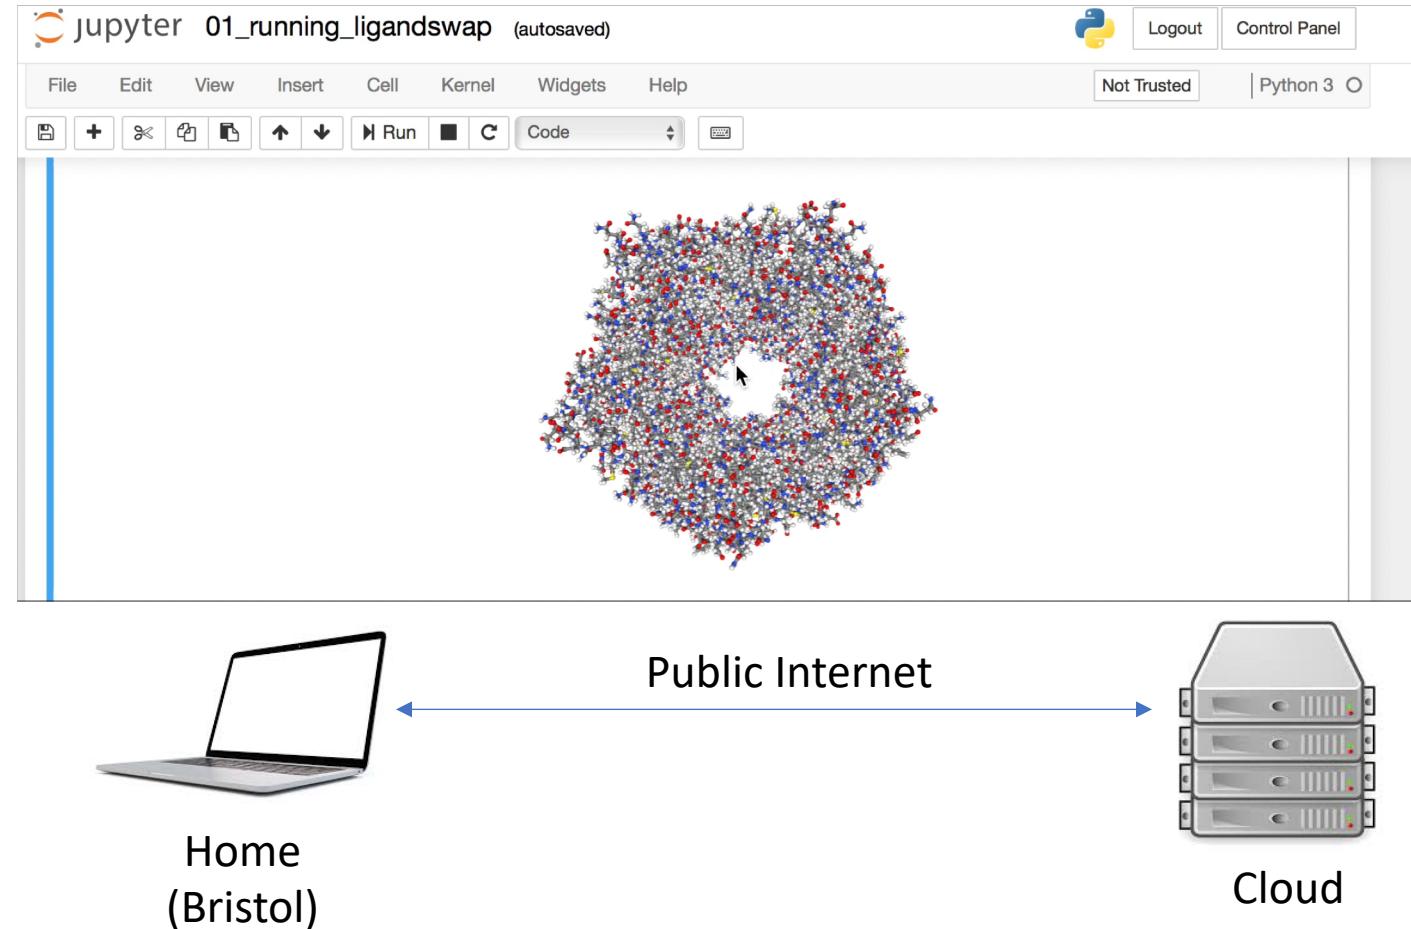


- **Jupyter notebook combines the;**
 - description of the experiment
 - code to run the experiment
 - code to analyse the results
 - graphs and 3D visualisations of the results
 - conclusions of the experiment
- **They contain everything needed to describe and reproduce the experiment**

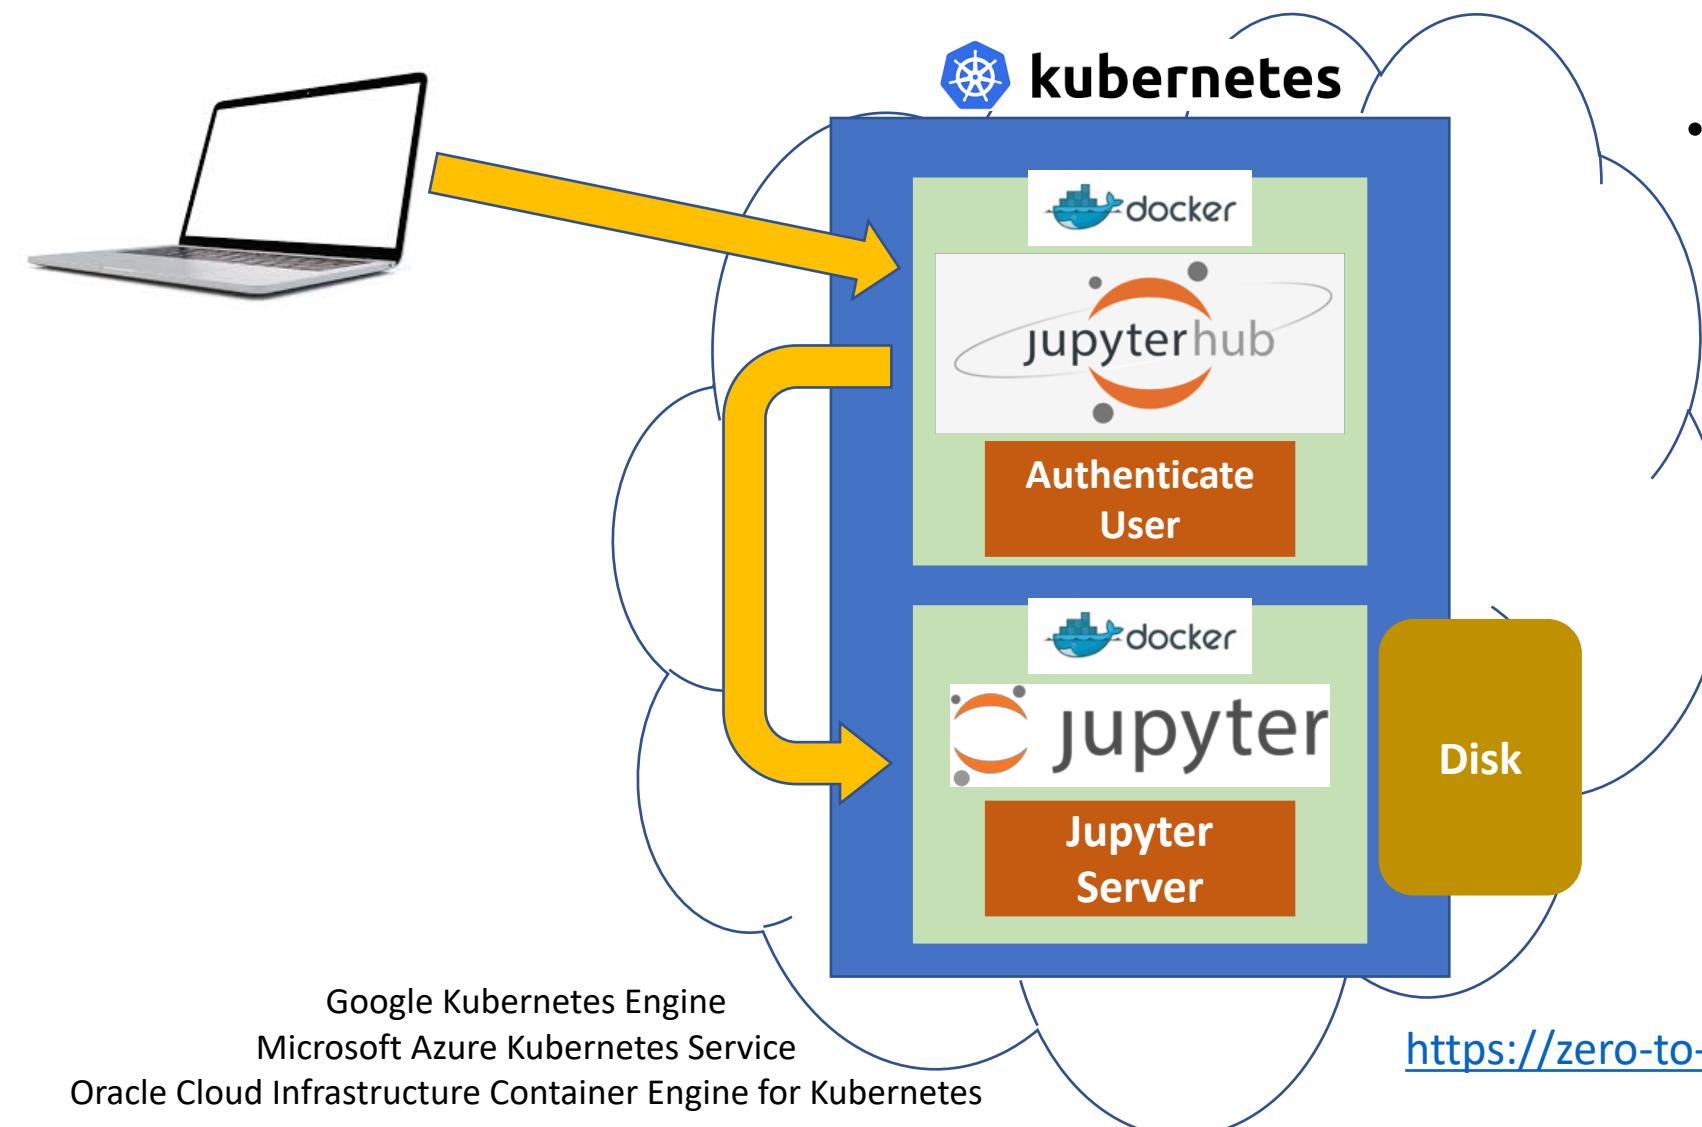
<https://jupyter.org> – <https://notebook.biosimspace.org/hub/tmplogin>

Notebooks are interactive papers

- All compute and data sits in the cloud. Only the information needed to render the data in the notebook is transmitted over the network
- Hugely useful for open and reproducible science
- Notebooks are, in effect, interactive scientific papers ☺



Kubernetes and JupyterHub



- **JupyterHub running on k8s**
 - Easy to use helm chart!
 - Great community and instructions
 - Works with lots (all?) cloud kubernetes services, or roll-your-own clusters

<https://zero-to-jupyterhub.readthedocs.io/en/stable/>

Simulations as Serverless Functions

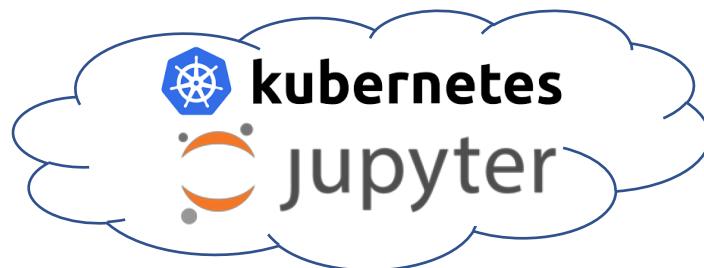
We now have everything that is needed to create a process object. To do so, run:

```
In [5]: process = BSS.MD.run(system, protocol)
```

```
In [6]: process.isRunning()
```

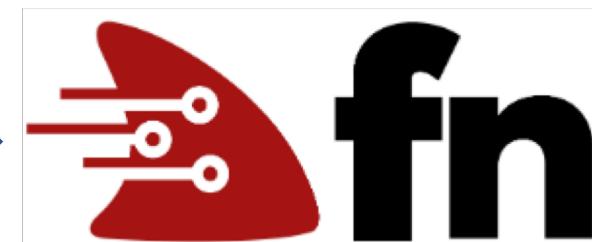
```
Out[6]: True
```

- The above line starts and runs a molecular dynamics simulation
- However, we cannot run this in the k8s pod, as the hardware is too tiny...
 - (...or else the k8s cluster would be too expensive)
- Instead we burst out to HPC hardware using a “serverless” function service



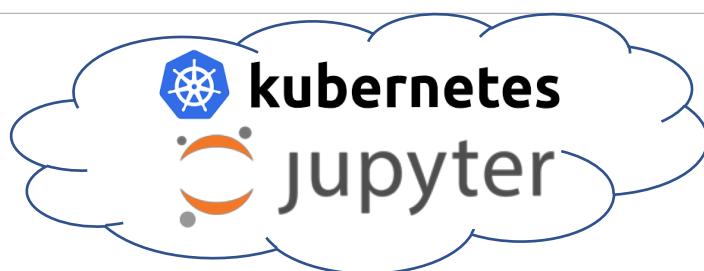
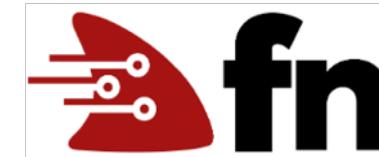
Auto-scaling 1 or 2 core VMs

`MD.run(system, protocol)`

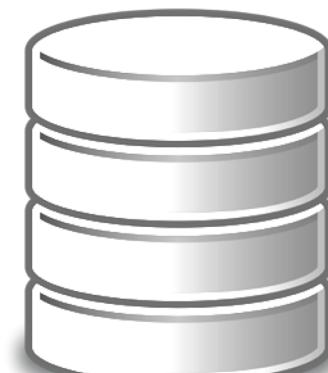


Simulations triggered by Fn Functions

```
In [ ]: import BioSimSpace as BSS  
  
In [ ]: system = BSS.IO.readMolecules(["amber/ala/ala.crd", "amber/ala/ala.top"])  
  
In [ ]: # Initialise a short equilibration protocol.  
protocol = BSS.Protocol.Equilibration(runtime=0.05*BSS.Units.Time.nanosecond,  
                                      temperature_start=0*BSS.Units.Temperature.kelvin,  
                                      temperature_end=300*BSS.Units.Temperature.kelvin,  
                                      restrain_backbone=True)  
  
In [ ]: process = BSS.MD.run(system, protocol)  
  
In [ ]: # Generate a plot of time vs temperature.  
plot1 = BSS.Notebook.plot(process.getTime(time_series=True),  
                         process.getTemperature(time_series=True))  
  
# Generate a plot of time vs energy.  
plot2 = BSS.Notebook.plot(process.getTime(time_series=True),  
                         process.getTotalEnergy(time_series=True))  
  
In [ ]: view = BSS.Notebook.View(process)  
view.system()
```

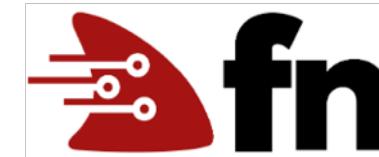


Auto-scaling 1 or 2 core VMs



Object Store

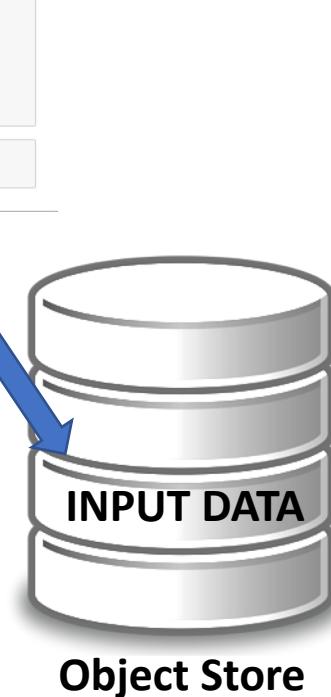
Simulations triggered by Fn Functions



```
In [ ]: import BioSimSpace as BSS  
  
In [ ]: system = BSS.IO.readMolecules(["amber/ala/ala.crd", "amber/ala/ala.top"])  
  
In [ ]: # Initialise a short equilibration protocol.  
protocol = BSS.Protocol.Equilibration(runtime=0.05*BSS.Units.Time.nanosecond,  
                                      temperature_start=0*BSS.Units.Temperature.kelvin,  
                                      temperature_end=300*BSS.Units.Temperature.kelvin,  
                                      restrain_backbone=True)  
  
In [ ]: process = BSS.MD.run(system, protocol)  
  
In [ ]: # Generate a plot of time vs temperature.  
plot1 = BSS.Notebook.plot(process.getTime(time_series=True),  
                           process.getTemperature(time_series=True))  
  
# Generate a plot of time vs energy.  
plot2 = BSS.Notebook.plot(process.getTime(time_series=True),  
                           process.getTotalEnergy(time_series=True))  
  
In [ ]: view = BSS.Notebook.View(process)  
view.system()
```



Auto-scaling 1 or 2 core VMs

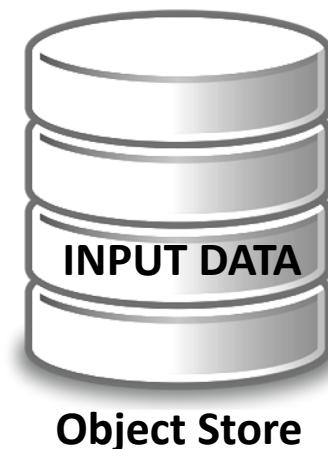


Simulations triggered by Fn Functions

```
In [ ]: import BioSimSpace as BSS  
  
In [ ]: system = BSS.IO.readMolecules(["amber/ala/ala.crd", "amber/ala/ala.top"])  
  
In [ ]: # Initialise a short equilibration protocol.  
protocol = BSS.Protocol.Equilibration(runtime=0.05*BSS.Units.Time.nanosecond,  
                                      temperature_start=0*BSS.Units.Temperature.kelvin,  
                                      temperature_end=300*BSS.Units.Temperature.kelvin,  
                                      restrain_backbone=True)  
  
In [ ]: process = BSS.MD.run(system, protocol)  
  
In [ ]: # Generate a plot of time vs temperature.  
plot1 = BSS.Notebook.plot(process.getTime(time_series=True),  
                           process.getTemperature(time_series=True))  
  
# Generate a plot of time vs energy.  
plot2 = BSS.Notebook.plot(process.getTime(time_series=True),  
                           process.getTotalEnergy(time_series=True))  
  
In [ ]: view = BSS.Notebook.View(process)  
view.system()
```

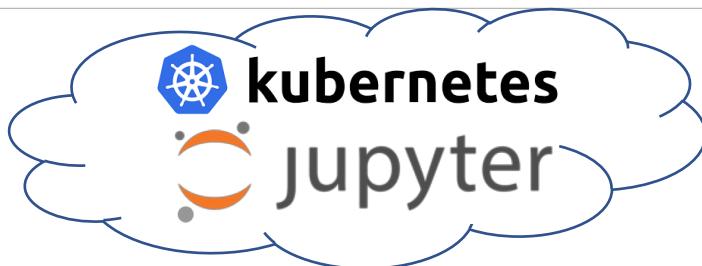


Auto-scaling 1 or 2 core VMs

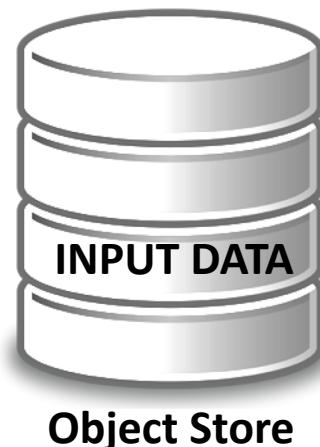


Simulations triggered by Fn Functions

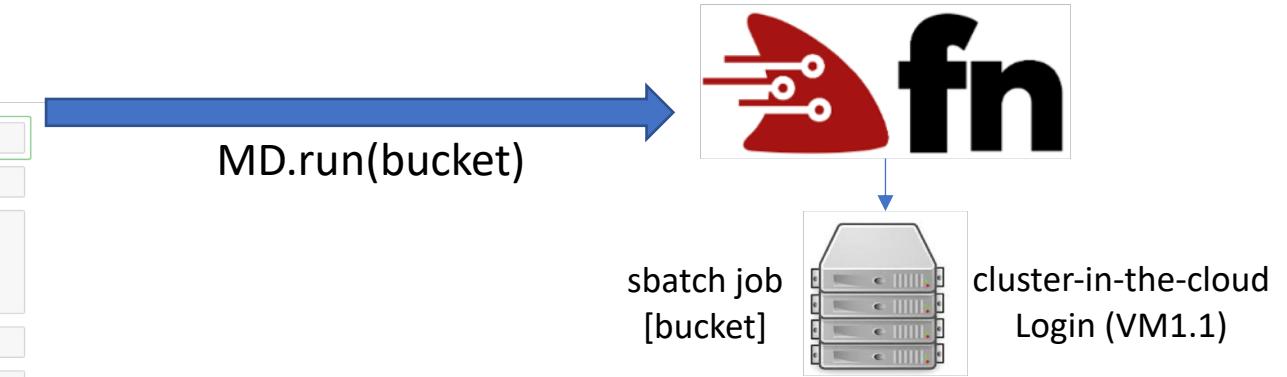
```
In [ ]: import BioSimSpace as BSS  
  
In [ ]: system = BSS.IO.readMolecules(["amber/ala/ala.crd", "amber/ala/ala.top"])  
  
In [ ]: # Initialise a short equilibration protocol.  
protocol = BSS.Protocol.Equilibration(runtime=0.05*BSS.Units.Time.nanosecond,  
                                      temperature_start=0*BSS.Units.Temperature.kelvin,  
                                      temperature_end=300*BSS.Units.Temperature.kelvin,  
                                      restrain_backbone=True)  
  
In [ ]: process = BSS.MD.run(system, protocol)  
  
In [ ]: # Generate a plot of time vs temperature.  
plot1 = BSS.Notebook.plot(process.getTime(time_series=True),  
                         process.getTemperature(time_series=True))  
  
# Generate a plot of time vs energy.  
plot2 = BSS.Notebook.plot(process.getTime(time_series=True),  
                         process.getTotalEnergy(time_series=True))  
  
In [ ]: view = BSS.Notebook.View(process)  
view.system()
```



Auto-scaling 1 or 2 core VMs

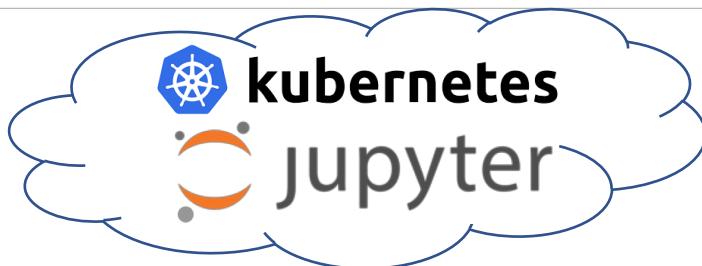


Object Store

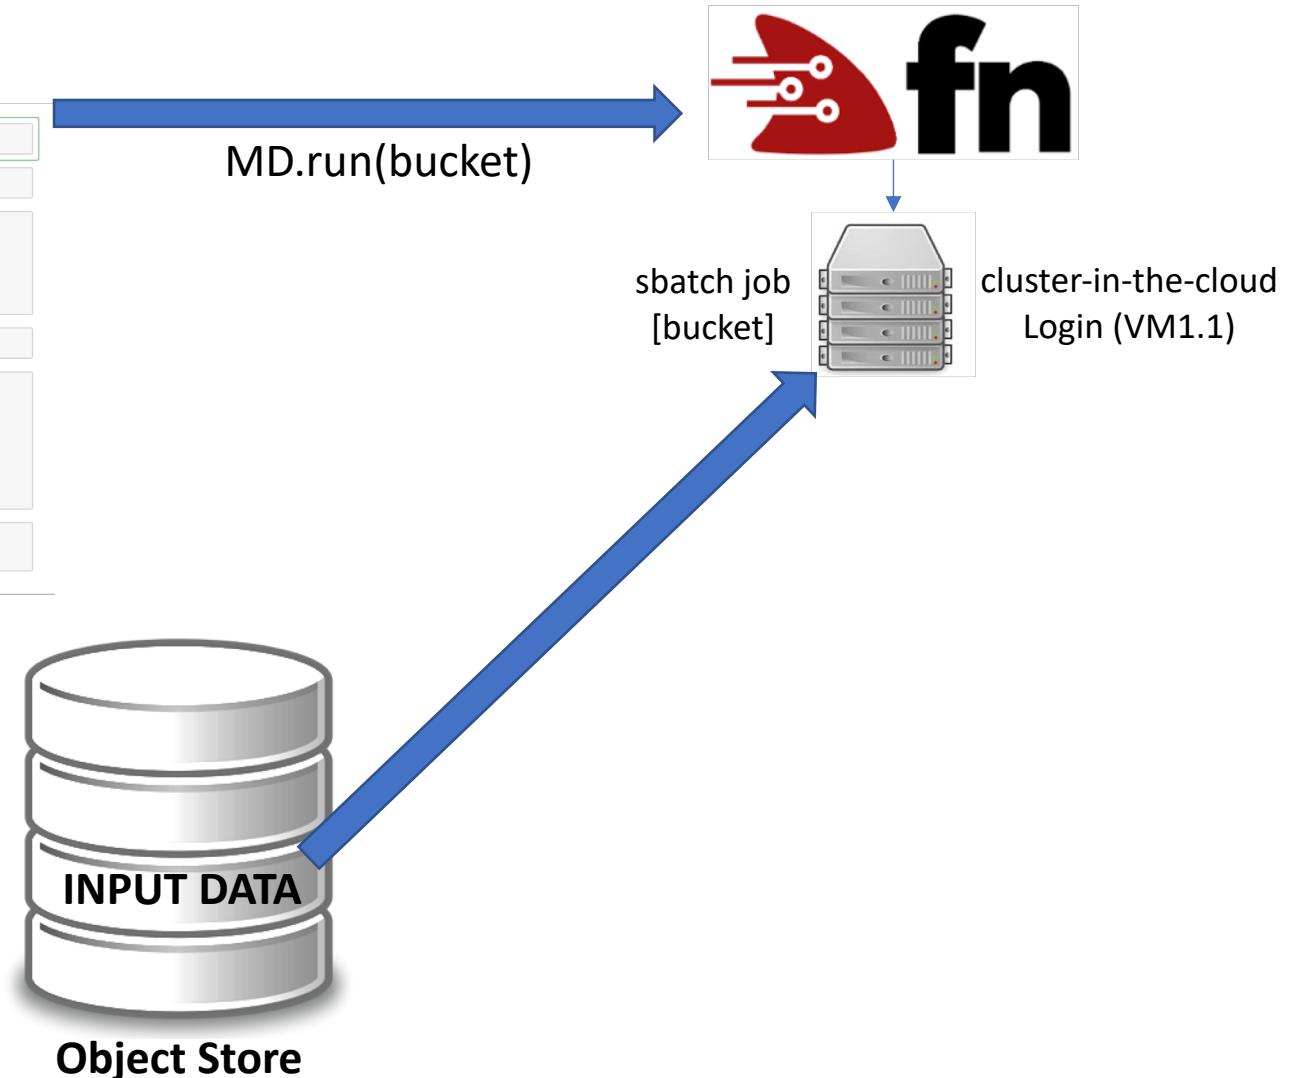


Simulations triggered by Fn Functions

```
In [ ]: import BioSimSpace as BSS  
  
In [ ]: system = BSS.IO.readMolecules(["amber/ala/ala.crd", "amber/ala/ala.top"])  
  
In [ ]: # Initialise a short equilibration protocol.  
protocol = BSS.Protocol.Equilibration(runtime=0.05*BSS.Units.Time.nanosecond,  
                                      temperature_start=0*BSS.Units.Temperature.kelvin,  
                                      temperature_end=300*BSS.Units.Temperature.kelvin,  
                                      restrain_backbone=True)  
  
In [ ]: process = BSS.MD.run(system, protocol)  
  
In [ ]: # Generate a plot of time vs temperature.  
plot1 = BSS.Notebook.plot(process.getTime(time_series=True),  
                         process.getTemperature(time_series=True))  
  
# Generate a plot of time vs energy.  
plot2 = BSS.Notebook.plot(process.getTime(time_series=True),  
                         process.getTotalEnergy(time_series=True))  
  
In [ ]: view = BSS.Notebook.View(process)  
view.system()
```

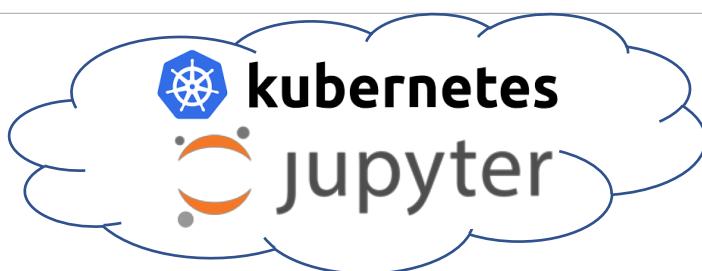


Auto-scaling 1 or 2 core VMs

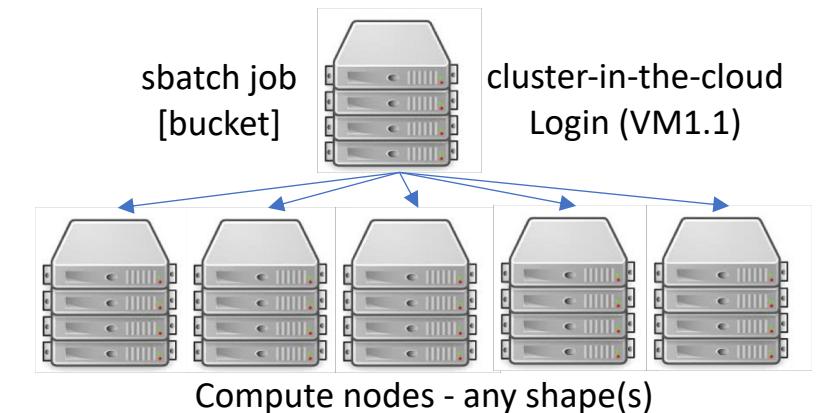
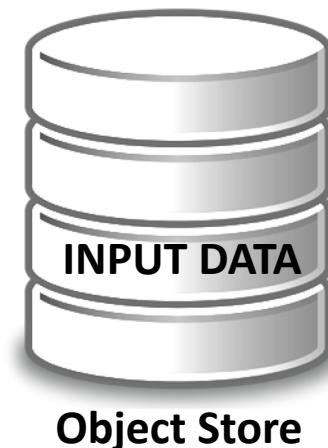


On-demand Cluster-in-the-Cloud

```
In [ ]: import BioSimSpace as BSS  
  
In [ ]: system = BSS.IO.readMolecules(["amber/ala/ala.crd", "amber/ala/ala.top"])  
  
In [ ]: # Initialise a short equilibration protocol.  
protocol = BSS.Protocol.Equilibration(runtime=0.05*BSS.Units.Time.nanosecond,  
                                      temperature_start=0*BSS.Units.Temperature.kelvin,  
                                      temperature_end=300*BSS.Units.Temperature.kelvin,  
                                      restrain_backbone=True)  
  
In [ ]: process = BSS.MD.run(system, protocol)  
  
In [ ]: # Generate a plot of time vs temperature.  
plot1 = BSS.Notebook.plot(process.getTime(time_series=True),  
                           process.getTemperature(time_series=True))  
  
# Generate a plot of time vs energy.  
plot2 = BSS.Notebook.plot(process.getTime(time_series=True),  
                           process.getTotalEnergy(time_series=True))  
  
In [ ]: view = BSS.Notebook.View(process)  
view.system()
```

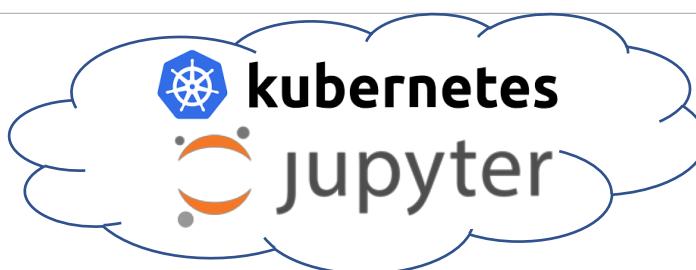


Auto-scaling 1 or 2 core VMs

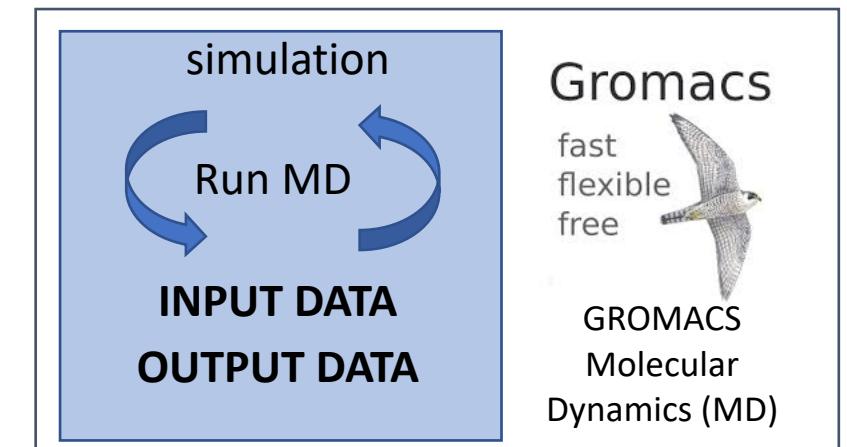
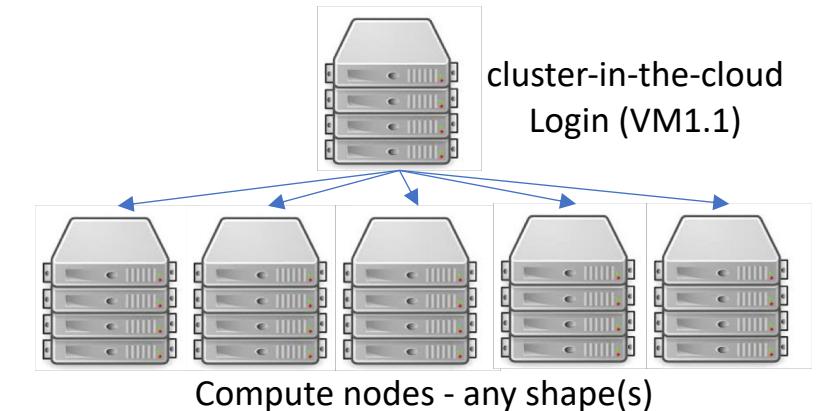
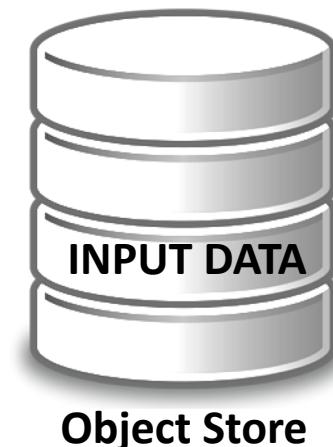


On-demand Cluster-in-the-Cloud

```
In [ ]: import BioSimSpace as BSS  
  
In [ ]: system = BSS.IO.readMolecules(["amber/ala/ala.crd", "amber/ala/ala.top"])  
  
In [ ]: # Initialise a short equilibration protocol.  
protocol = BSS.Protocol.Equilibration(runtime=0.05*BSS.Units.Time.nanosecond,  
                                      temperature_start=0*BSS.Units.Temperature.kelvin,  
                                      temperature_end=300*BSS.Units.Temperature.kelvin,  
                                      restrain_backbone=True)  
  
In [ ]: process = BSS.MD.run(system, protocol)  
  
In [ ]: # Generate a plot of time vs temperature.  
plot1 = BSS.Notebook.plot(process.getTime(time_series=True),  
                           process.getTemperature(time_series=True))  
  
# Generate a plot of time vs energy.  
plot2 = BSS.Notebook.plot(process.getTime(time_series=True),  
                           process.getTotalEnergy(time_series=True))  
  
In [ ]: view = BSS.Notebook.View(process)  
view.system()
```

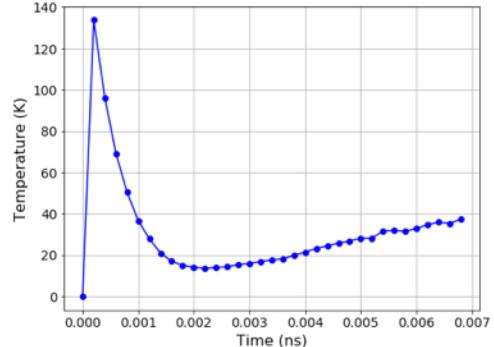


Auto-scaling 1 or 2 core VMs

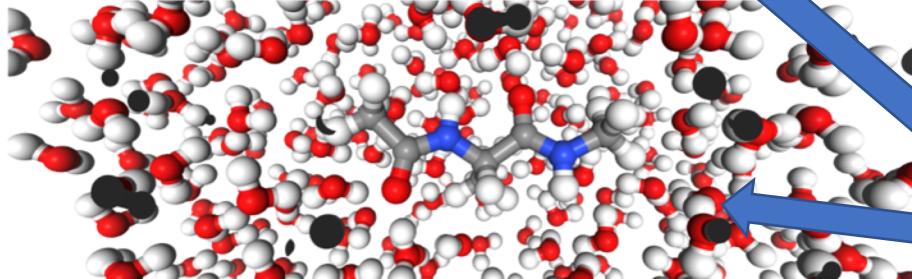


Serverless + Object Store 😊

```
In [4]: process = BSS.MD.run(system, protocol)  
  
In [6]: # Generate a plot of time vs temperature.  
plot1 = BSS.Notebook.plot(process.getTime(time_series=True),  
    process.getTemperature(time_series=True))
```

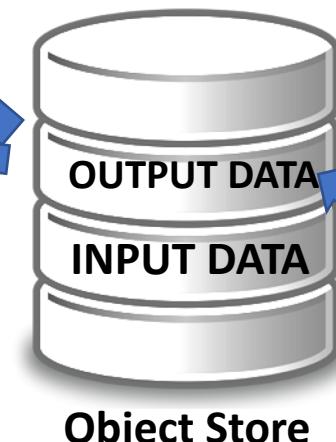


```
In [7]: view = BSS.Notebook.View(process)  
view.system()
```

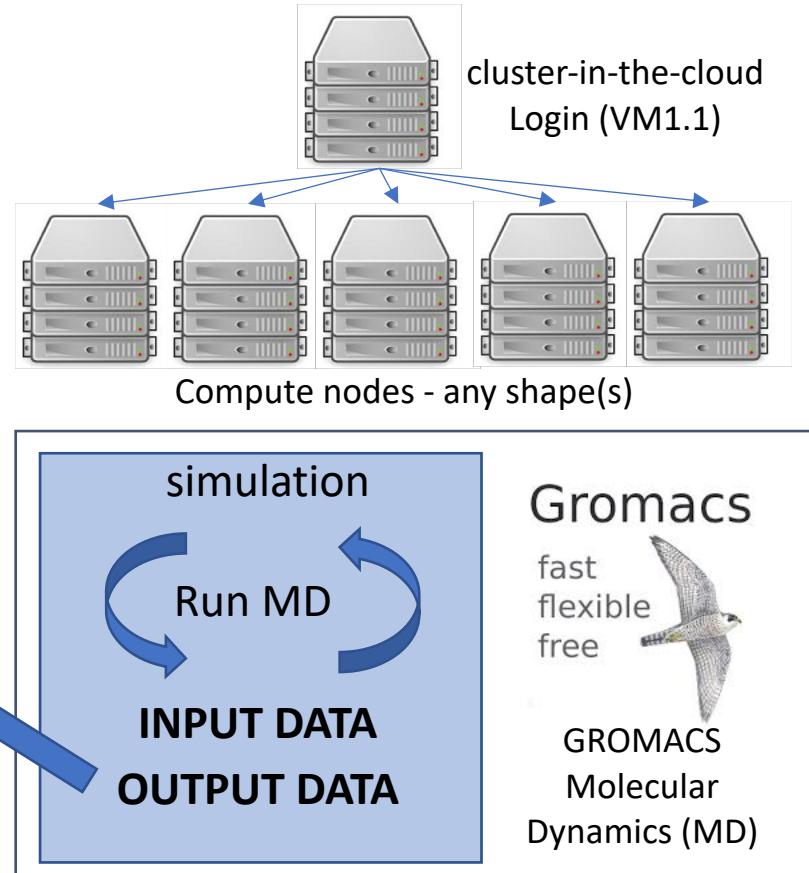


Auto-scaling 1 or 2 core VMs

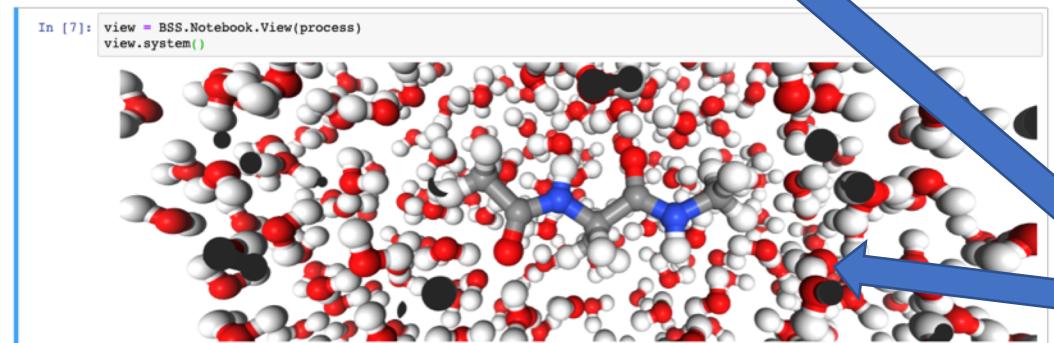
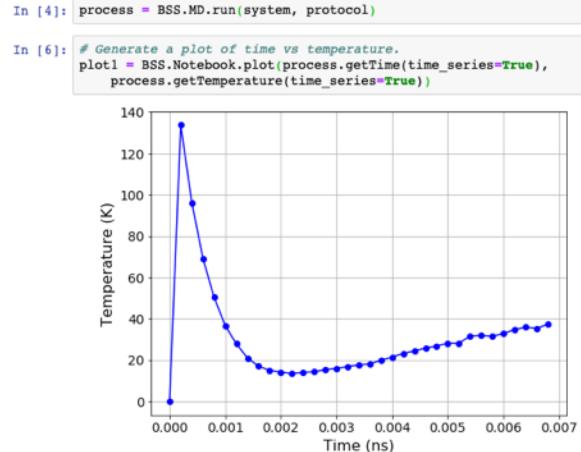
Live analysis by querying data as it arrives in the object store



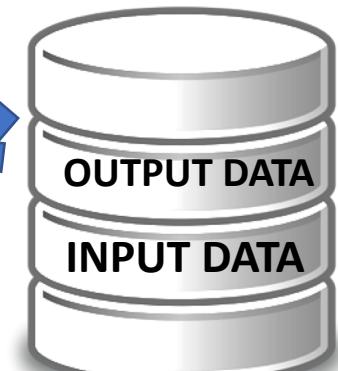
Stream output to object store



Serverless + Object Store 😊



Live analysis by querying data as it arrives in the object store

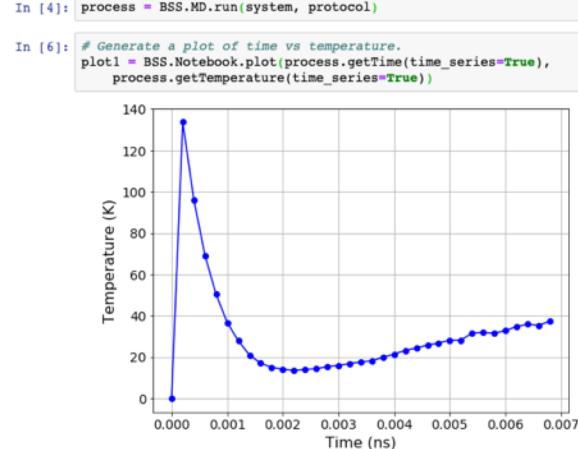


cluster-in-the-cloud
Login (VM1.1)

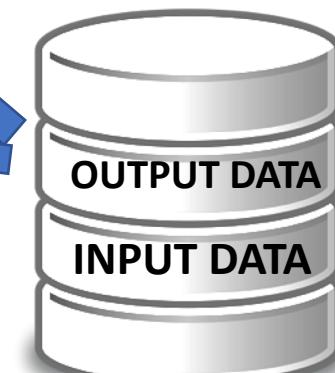
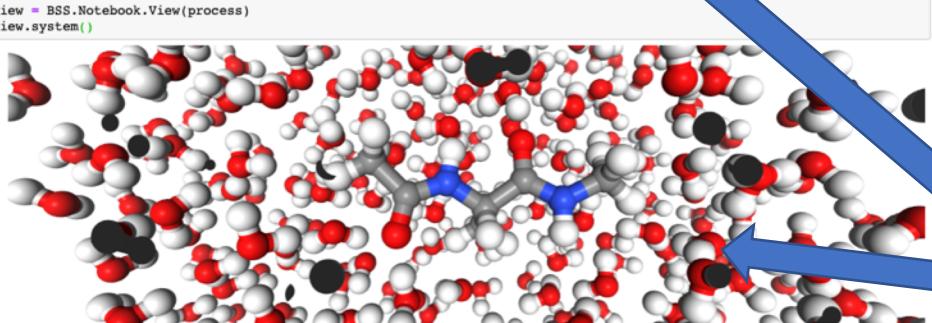


Auto-scaling 1 or 2 core VMs

Serverless + Object Store 😊



Live analysis by querying data as it arrives in the object store



Auto-scaling 1 or 2 core VMs

Anyone can run simulations...!

- Simple framework that allows ANYONE to run HPC simulations by calling the Fn function via a public URL!

Anyone can run simulations...!

- Simple framework that allows ANYONE to run HPC simulations by calling the Fn function via a public URL!



That could get expensive...!

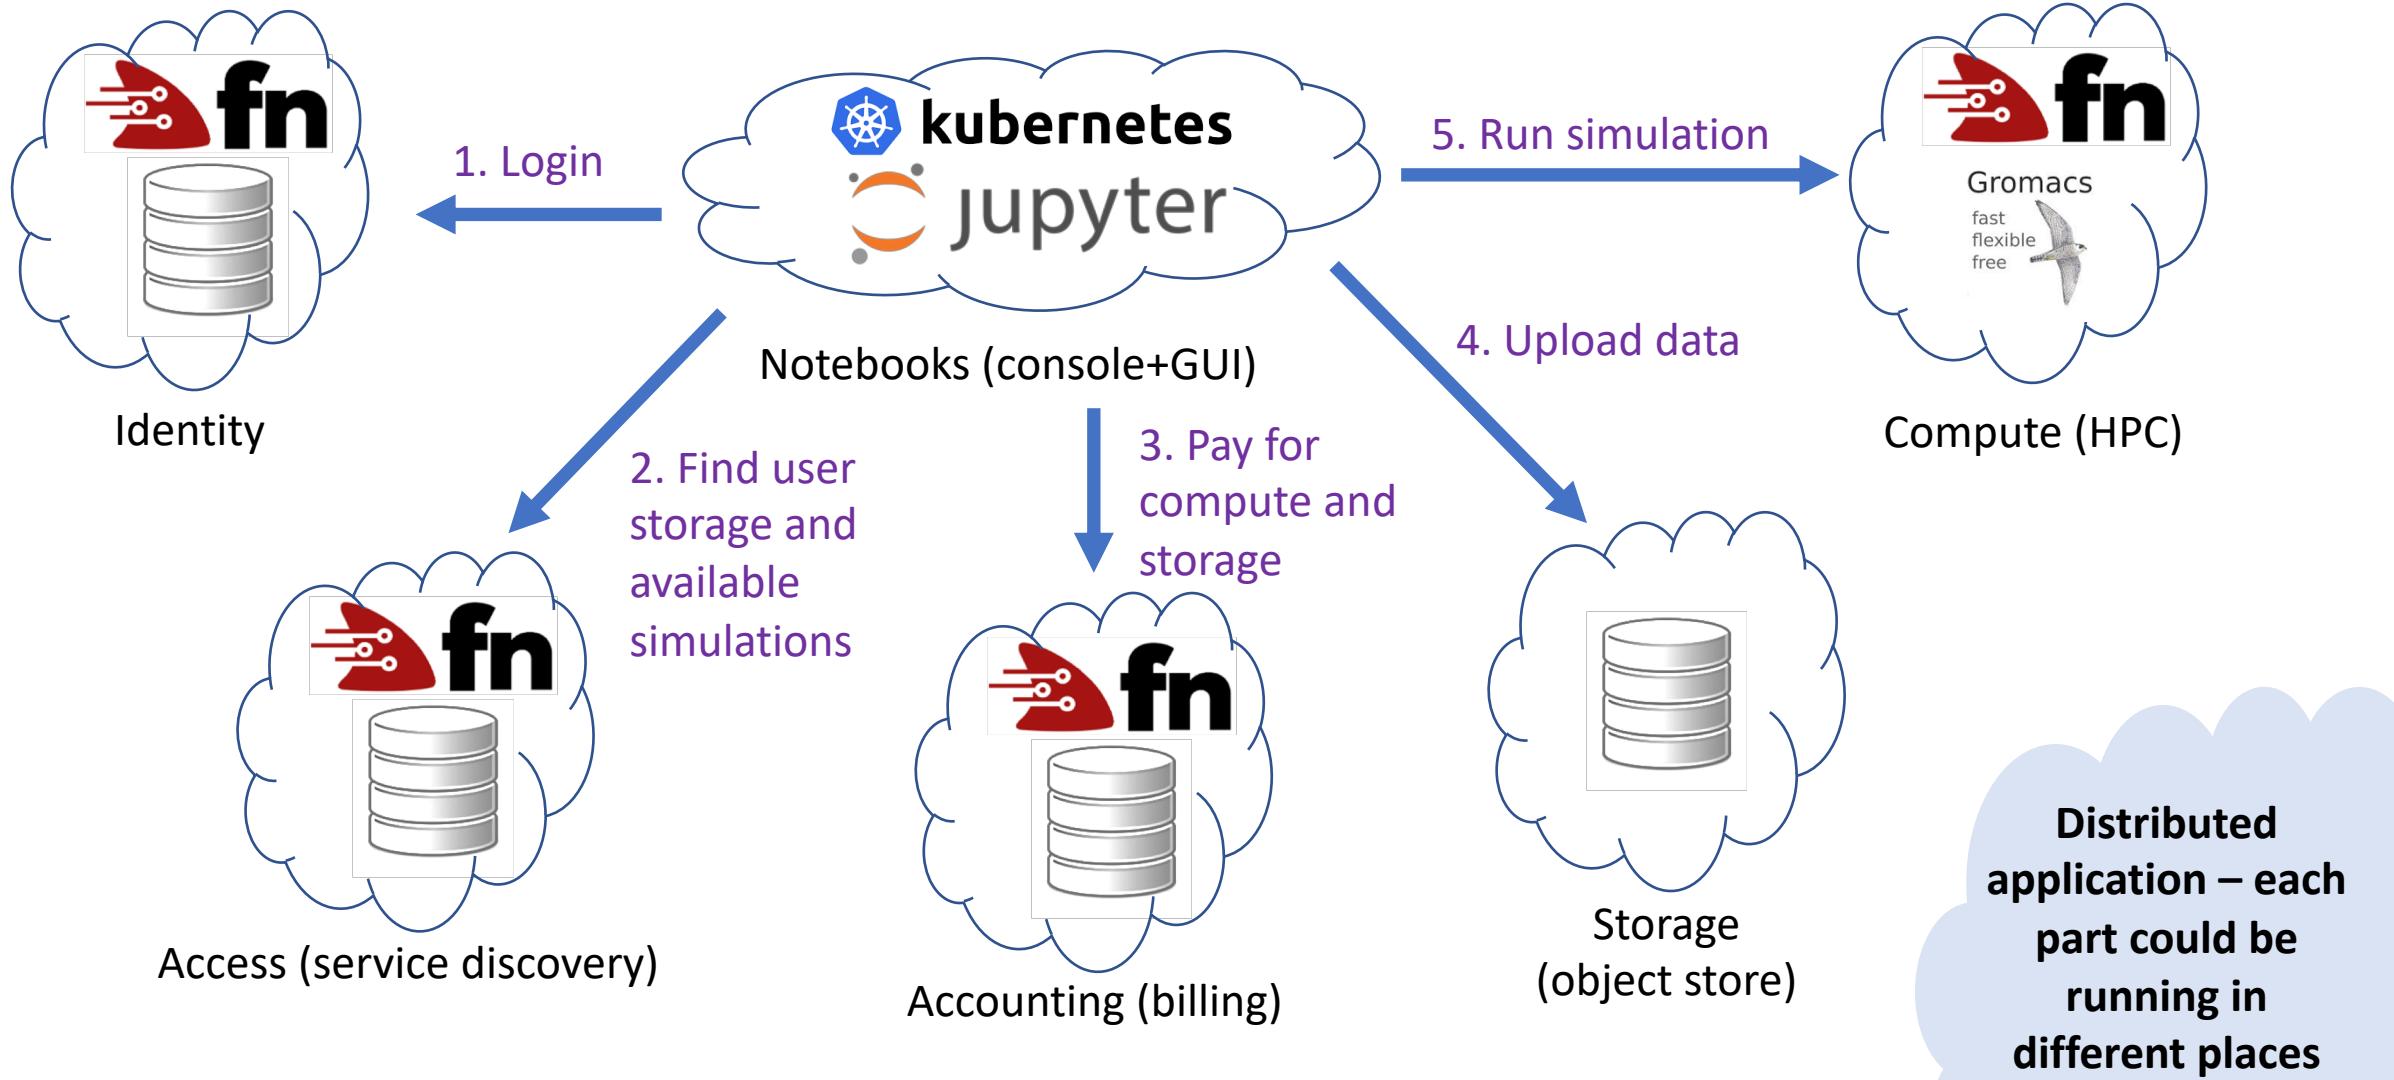
Looks like we need some user authentication, access control and accounting...

Identity (authorisation)

```
In [ ]: from Acquire.Client import User  
  
In [ ]: user = User("chryswoods")  
  
In [ ]: (url,qrcode) = user.request_login()  
  
In [ ]: qrcode  
  
In [ ]: from IPython.core.display import HTML  
print(url)  
HTML("<a href='%s'>Login here</a>" % url)  
  
In [ ]: user.wait_for_login()  
  
In [ ]: user.is_logged_in()  
  
In [ ]: user.logout()  
  
In [ ]:
```

- Built identity (authorisation) service on top of Fn serverless and object store for state
- “request_login” call from the notebook calls “request_login” serverless function. This looks up user details from object store and returns a unique login URL
- Login page also served as html from an Fn function
- Notebook can wait for the login to complete, and uses security tokens to authenticate with simulation function service
- <https://github.com/chryswoods/acquire>

Identity, Access, Accounting, Storage, Compute



Acquire = Cluster-in-the-Cloud + Accounting + Data Management

- Aim to allow **upfront charging** of cloud compute and storage costs on a “**per simulation**” basis
- Replaces current low-level charging based on VMs, network, disk etc.
- **Acquire will estimate compute and storage requirement of a simulation and will present an up-front guaranteed cost to run the calculation**
- Users can set **daily caps** and **maximum runtimes**
- **Best resource that fits the caps will be automatically chosen**

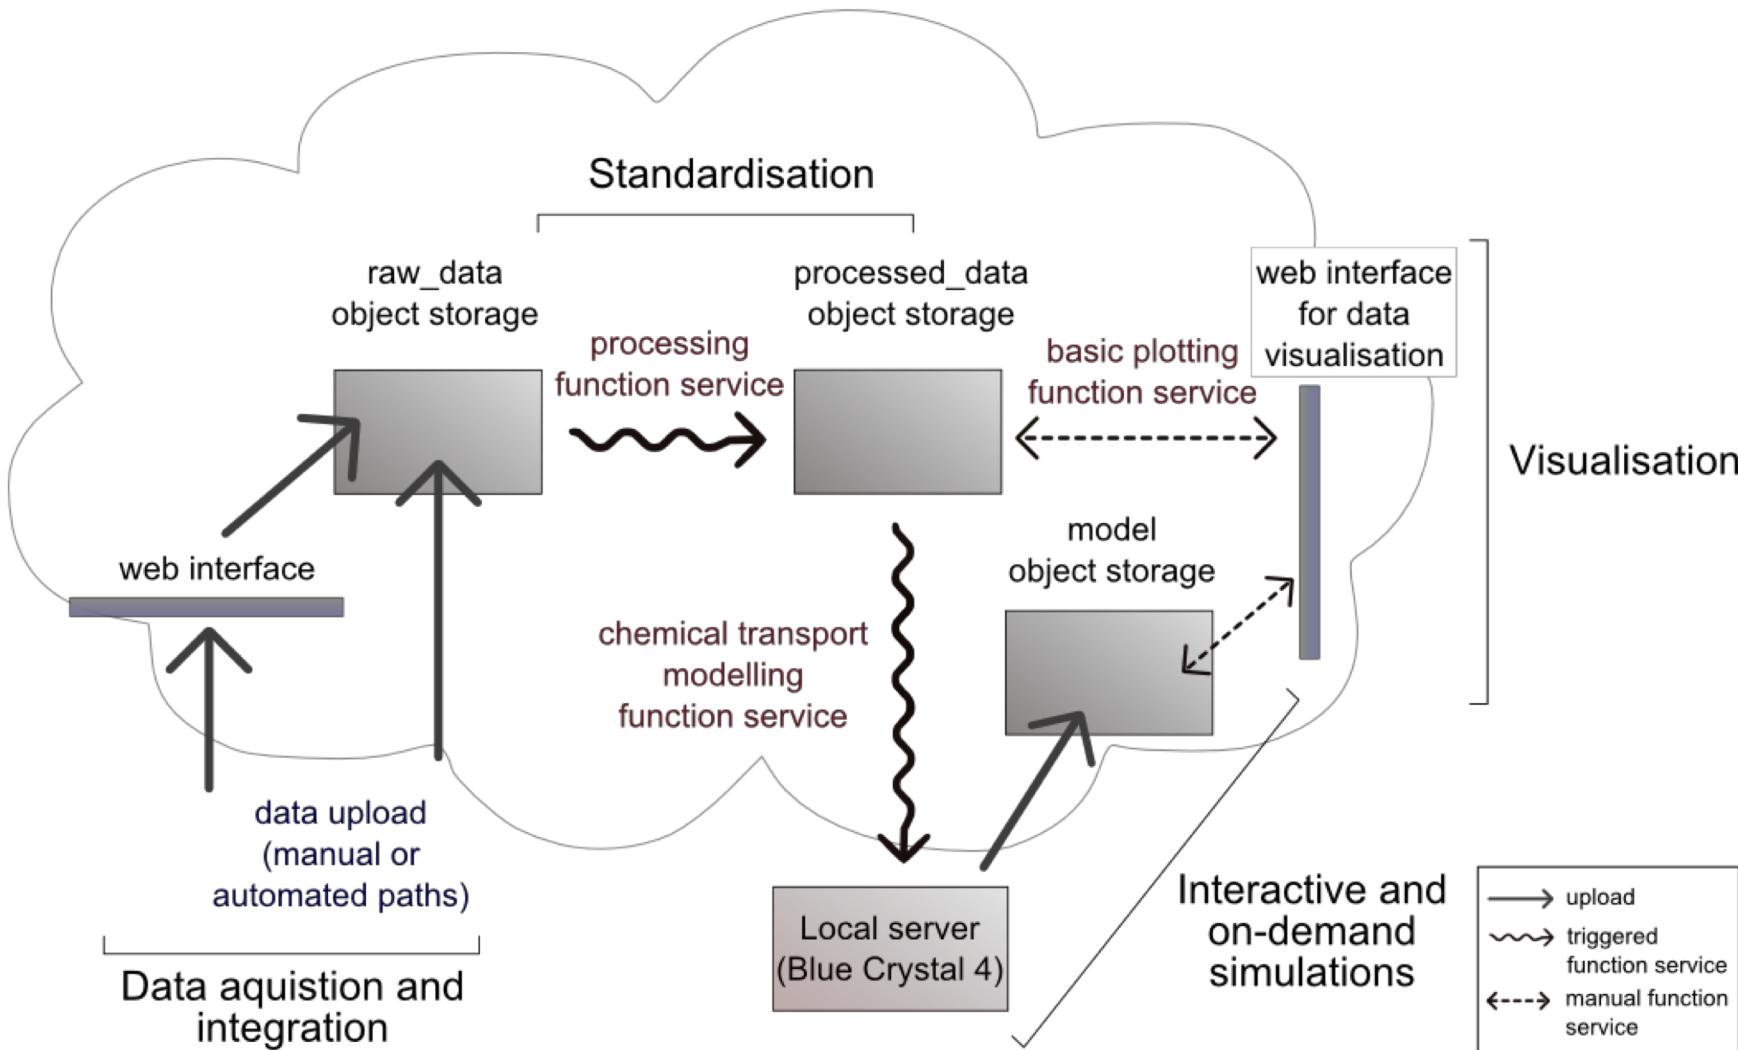
Acquire = Cluster-in-the-Cloud + Accounting + Data Management

- Outputs will be **read-only** and up-front **cost covers one year's storage in the object store**
- Object store key can become a **DOI** that allows them to be **accessed from other scripts, or published and accessed by others**
- Planning a web console to allow researchers to manage outputs, e.g. **control access permissions, delete the output** (receive a pro-rata storage refund), pay off extra years storage, or pay a **one-off charge for the output to be archived** (15 years)

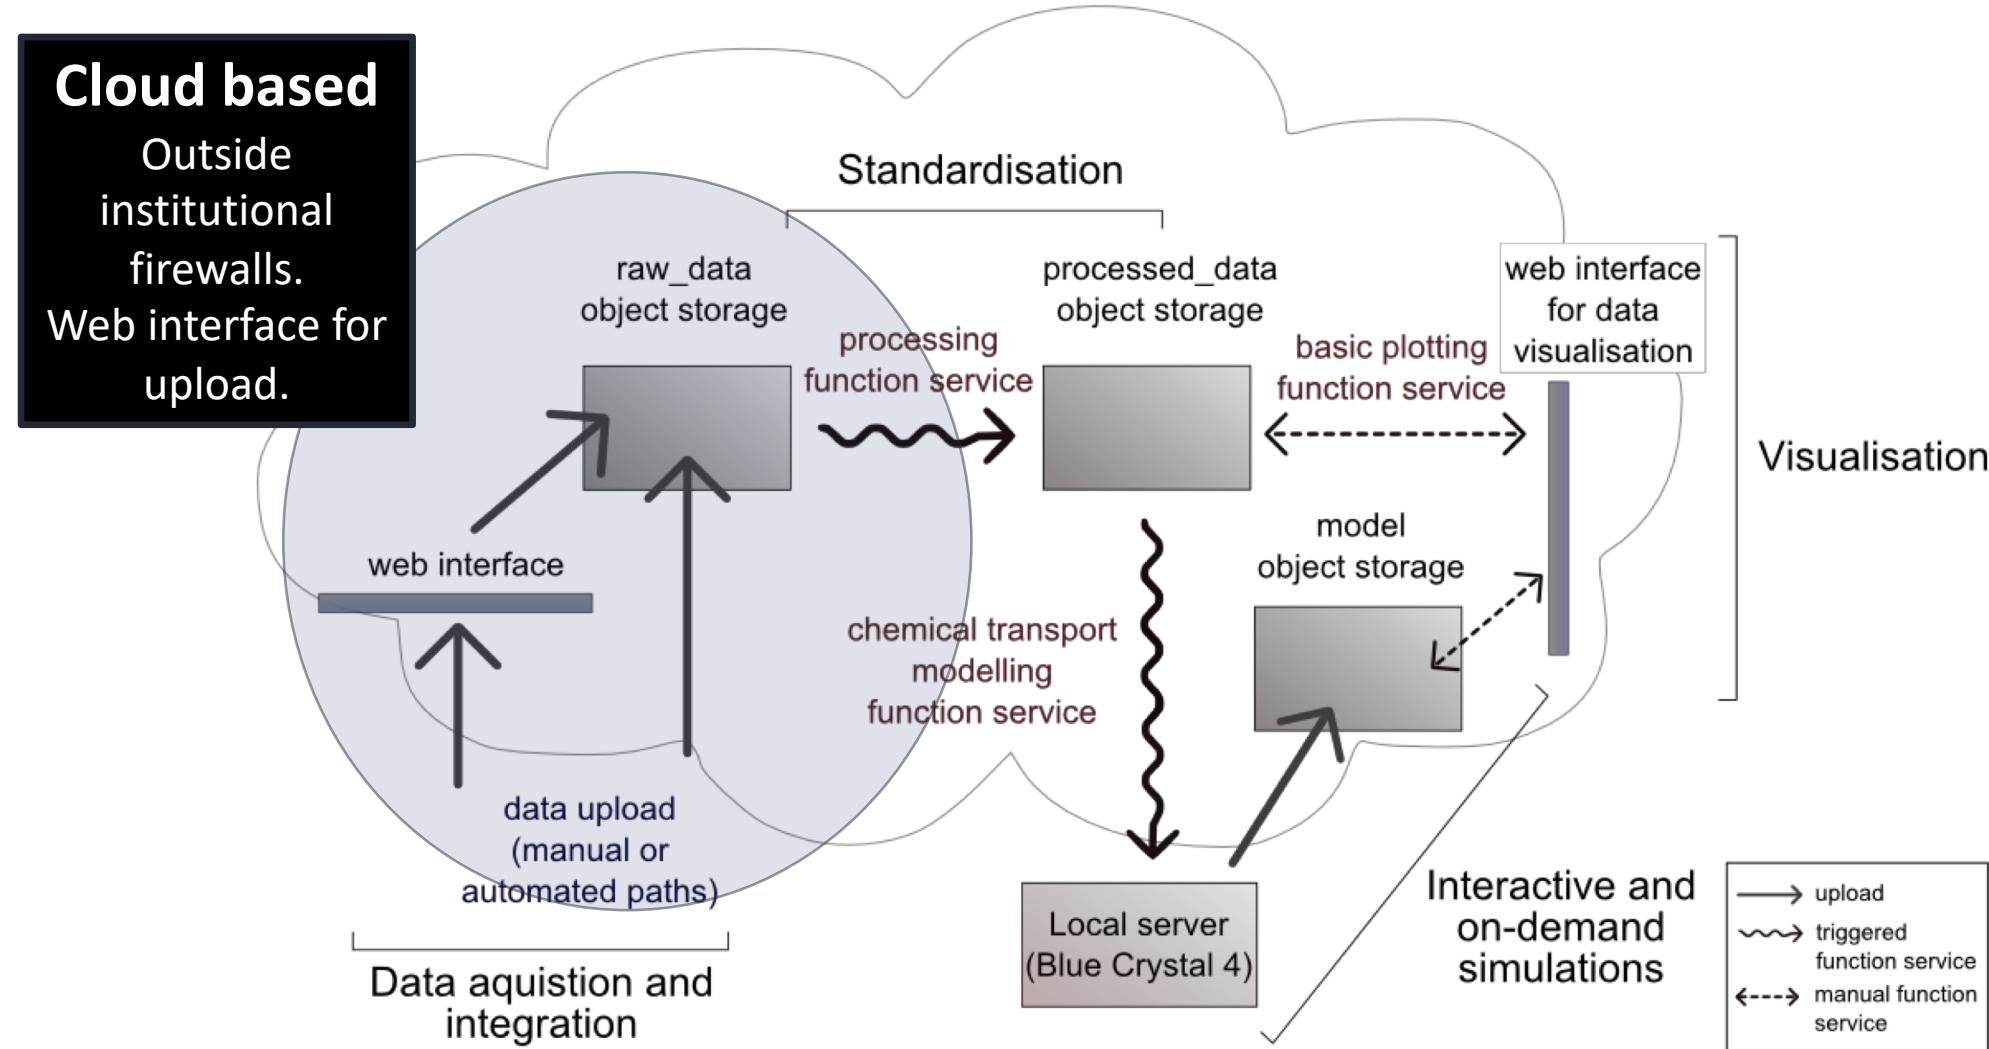
Acquire = Cluster-in-the-Cloud + Accounting + Data Management

- System will **support Jupyter-notebook type users, and also “traditional” users** running normal HPC workloads via a cluster-in-the-cloud interface.
- It could **scale from a single researcher** on one project using notebooks...
- ...**through universities managing 1000’s users** on “infinitely” scaling “infinitely” heterogeneous HPC clusters in the cloud...
- ...**to researchers across many universities and companies** sharing and collaborating on data running calculations, analyses, publishing interactive papers and **performing modern digital science**

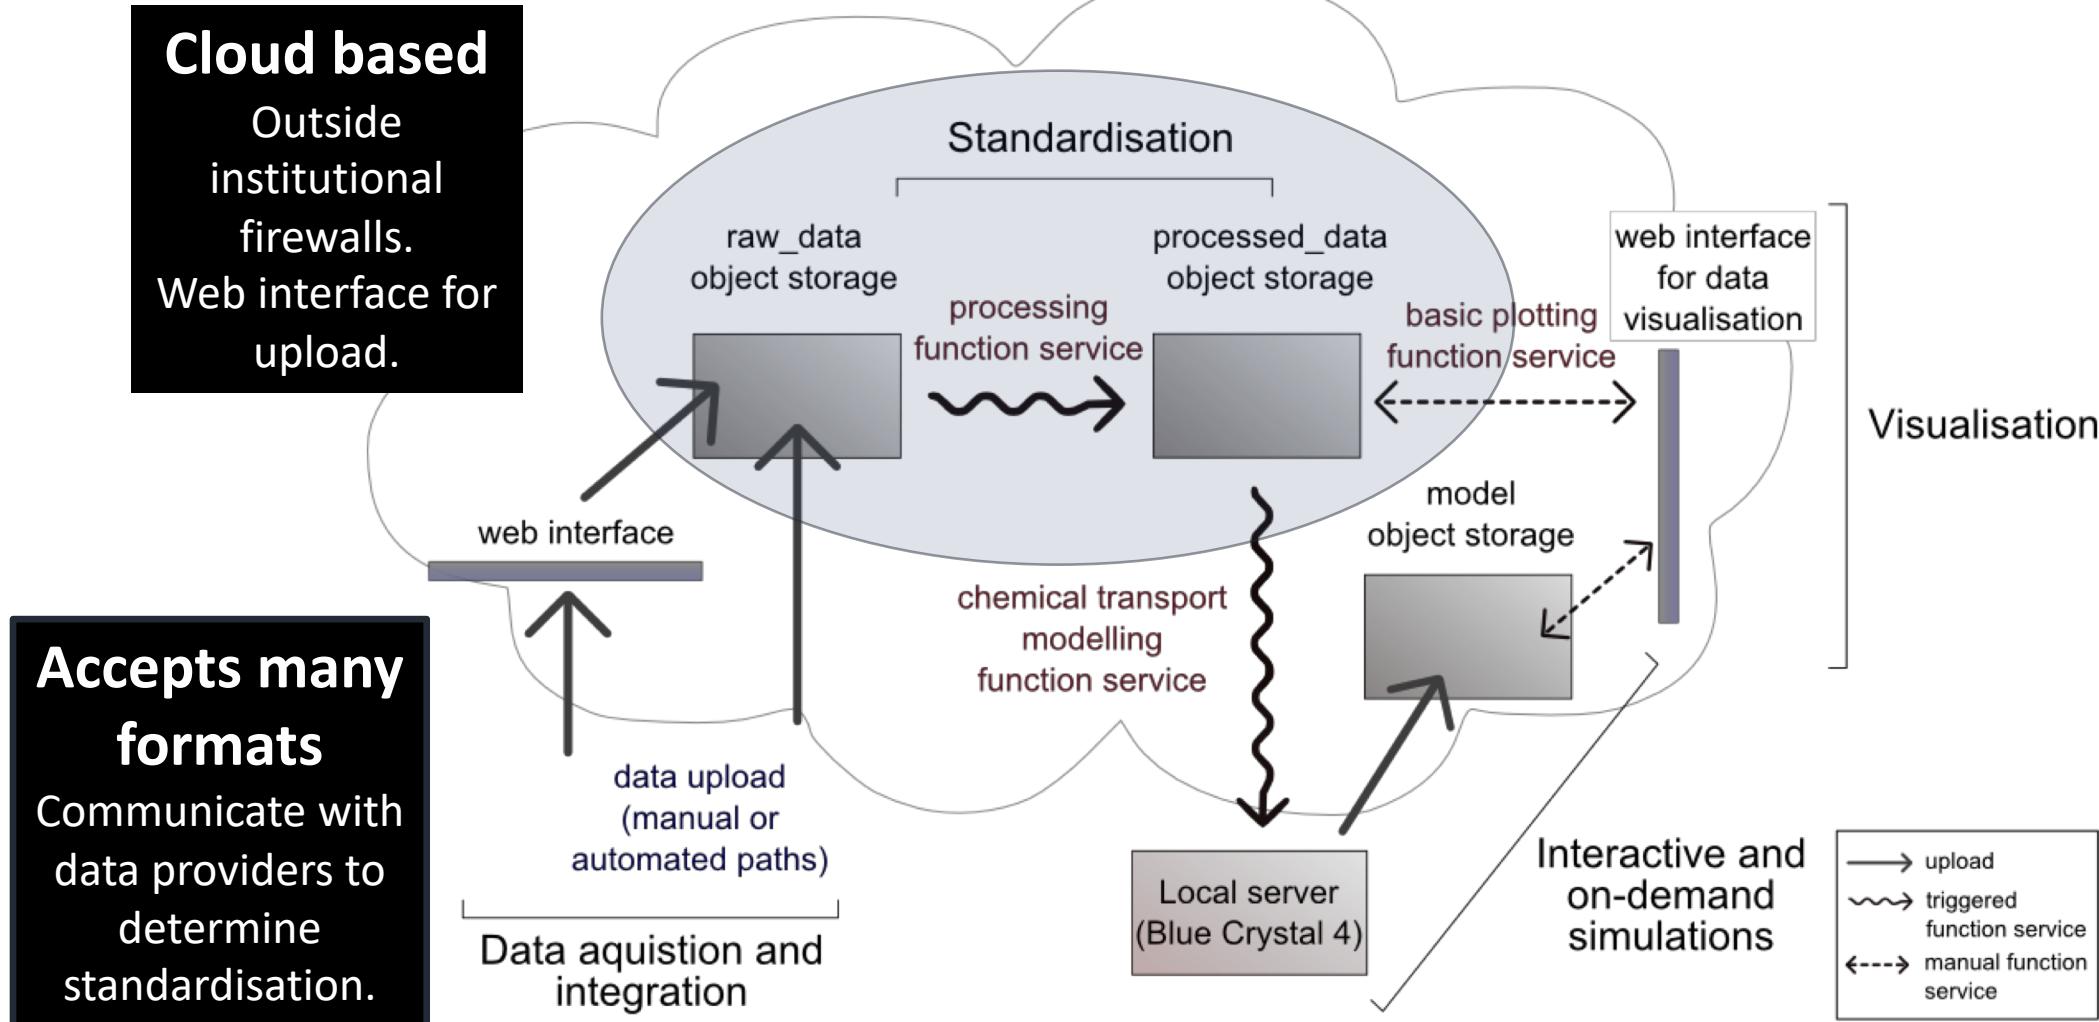
HUGS (Hub for UK Greenhouse Gas Science)



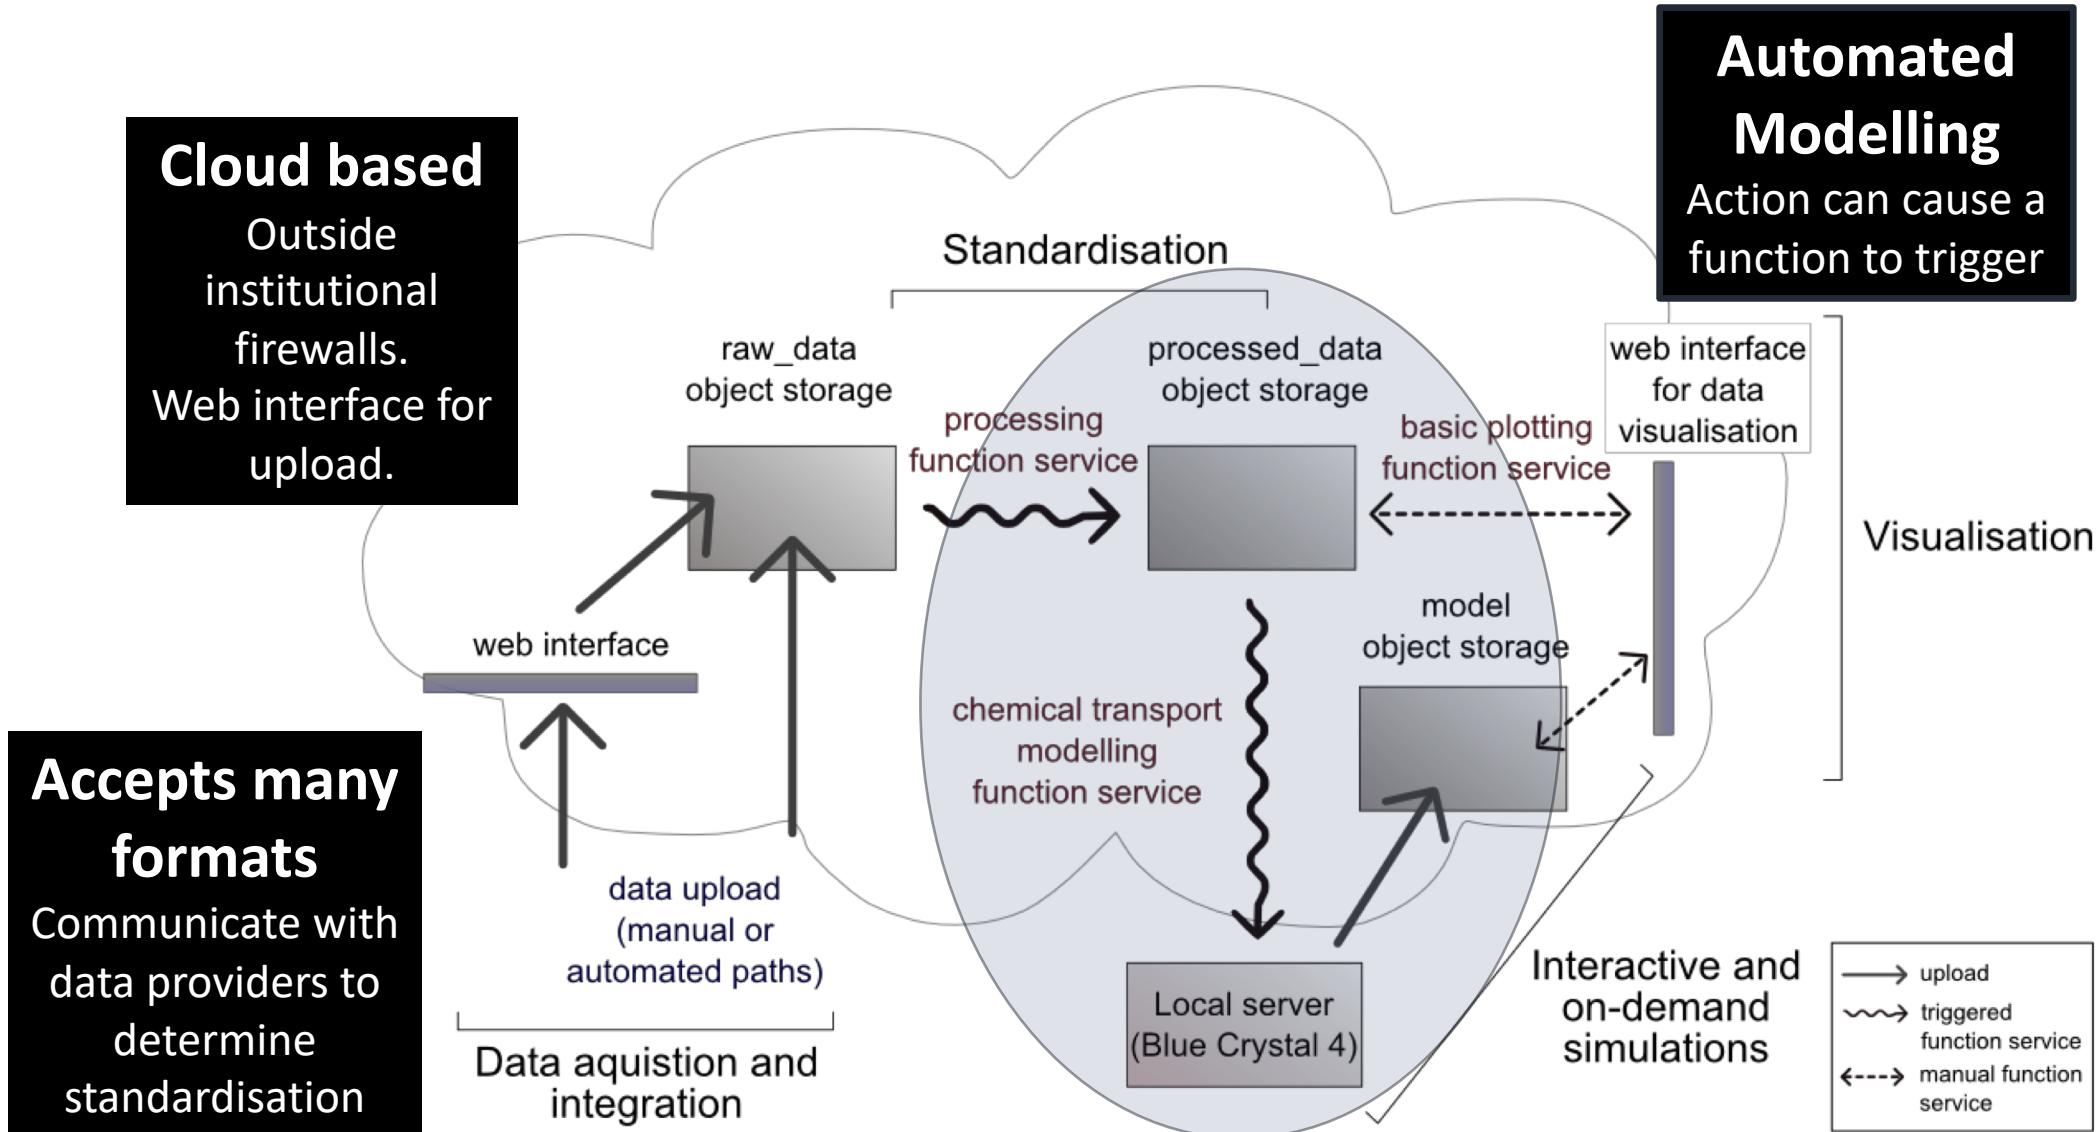
HUGS (Hub for UK Greenhouse Gas Science)



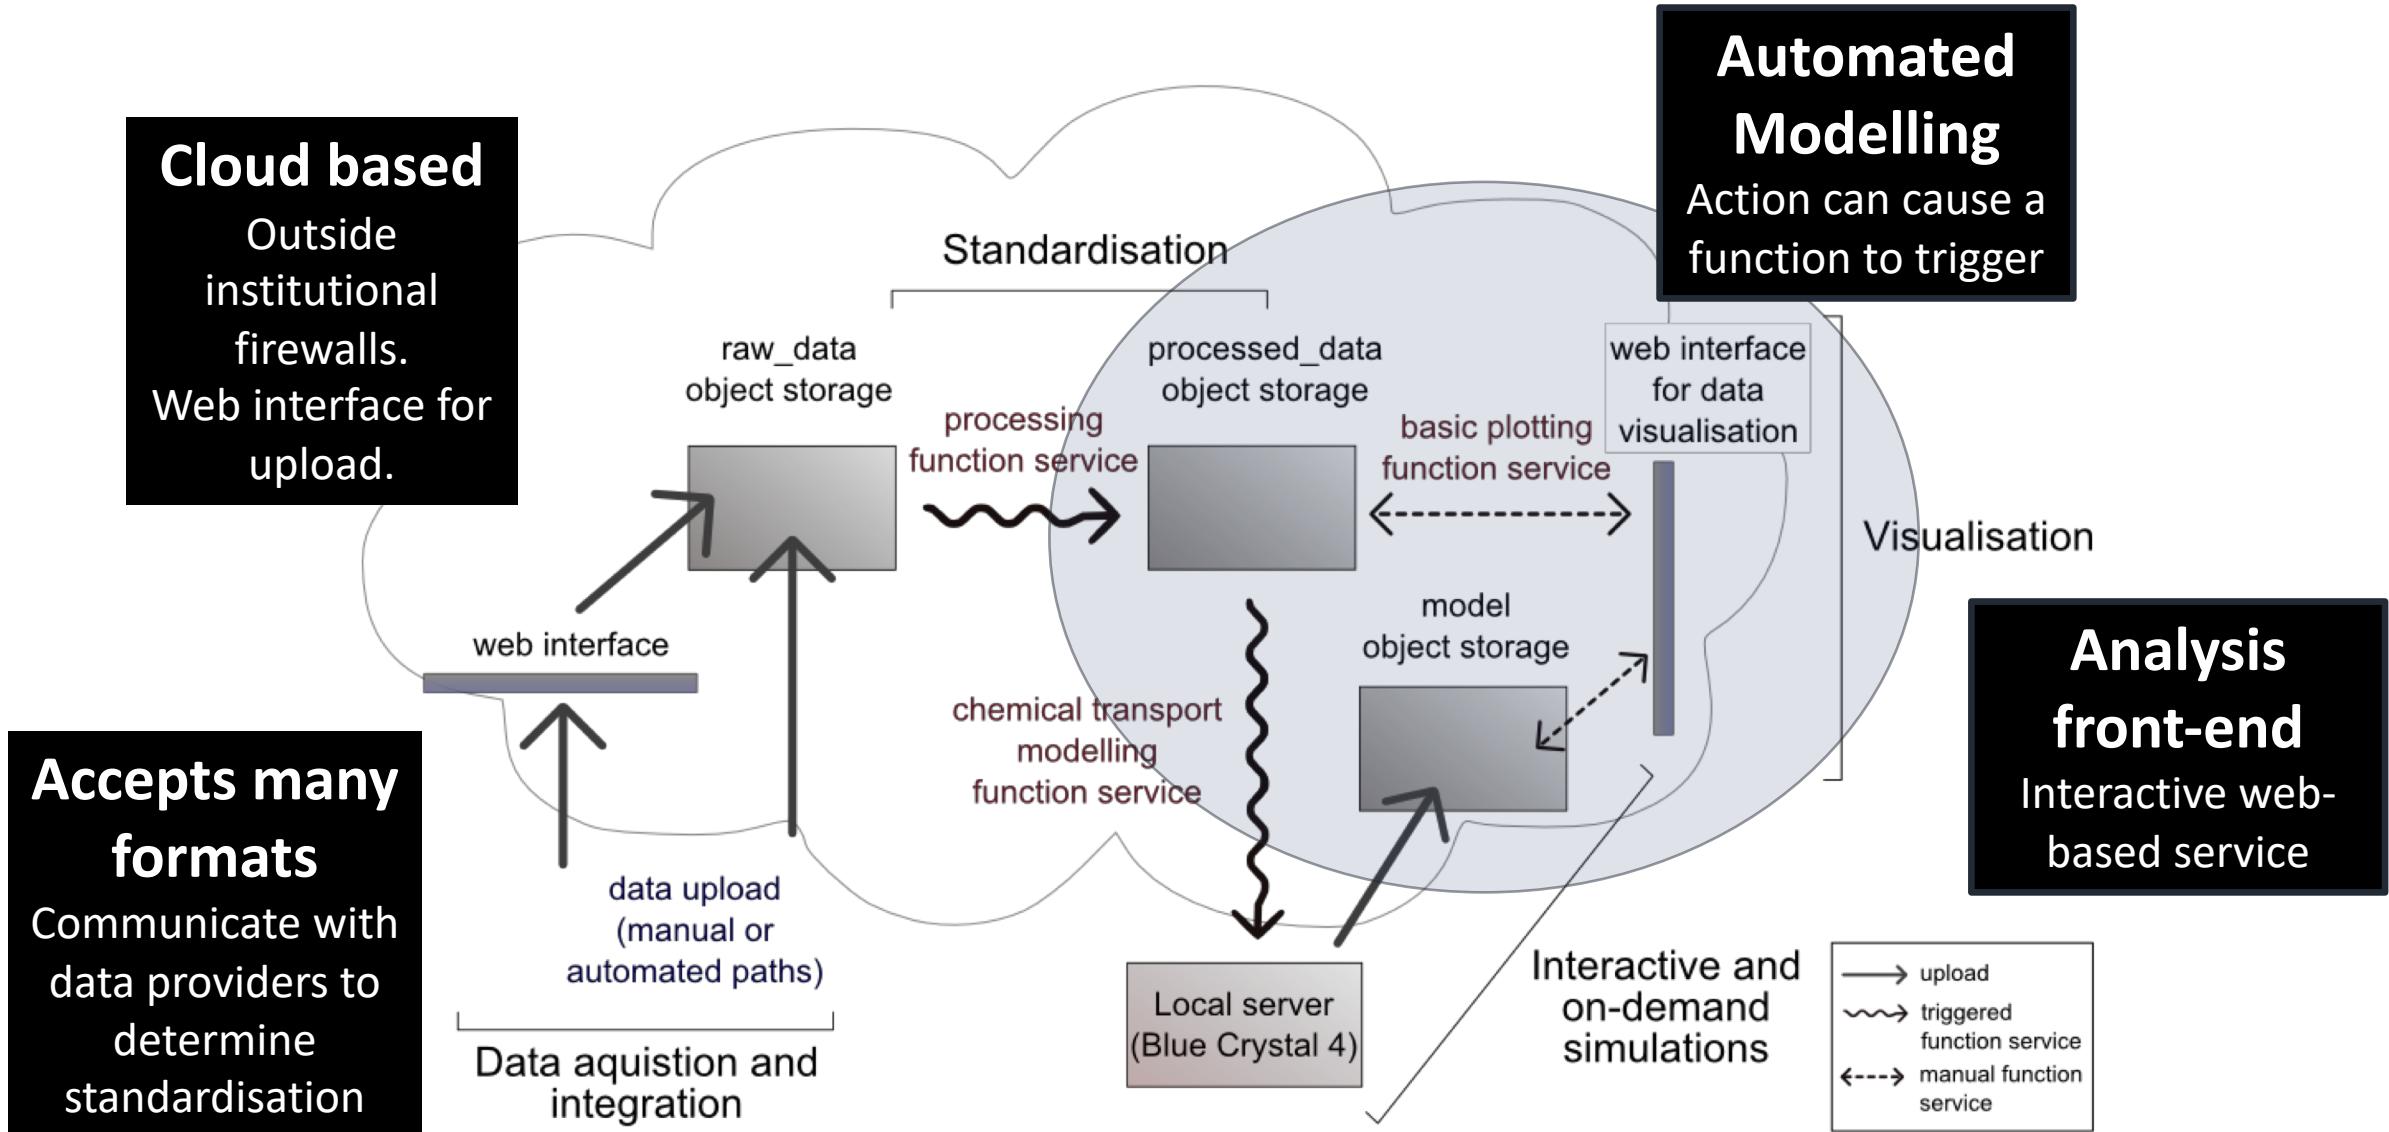
HUGS (Hub for UK Greenhouse Gas Science)

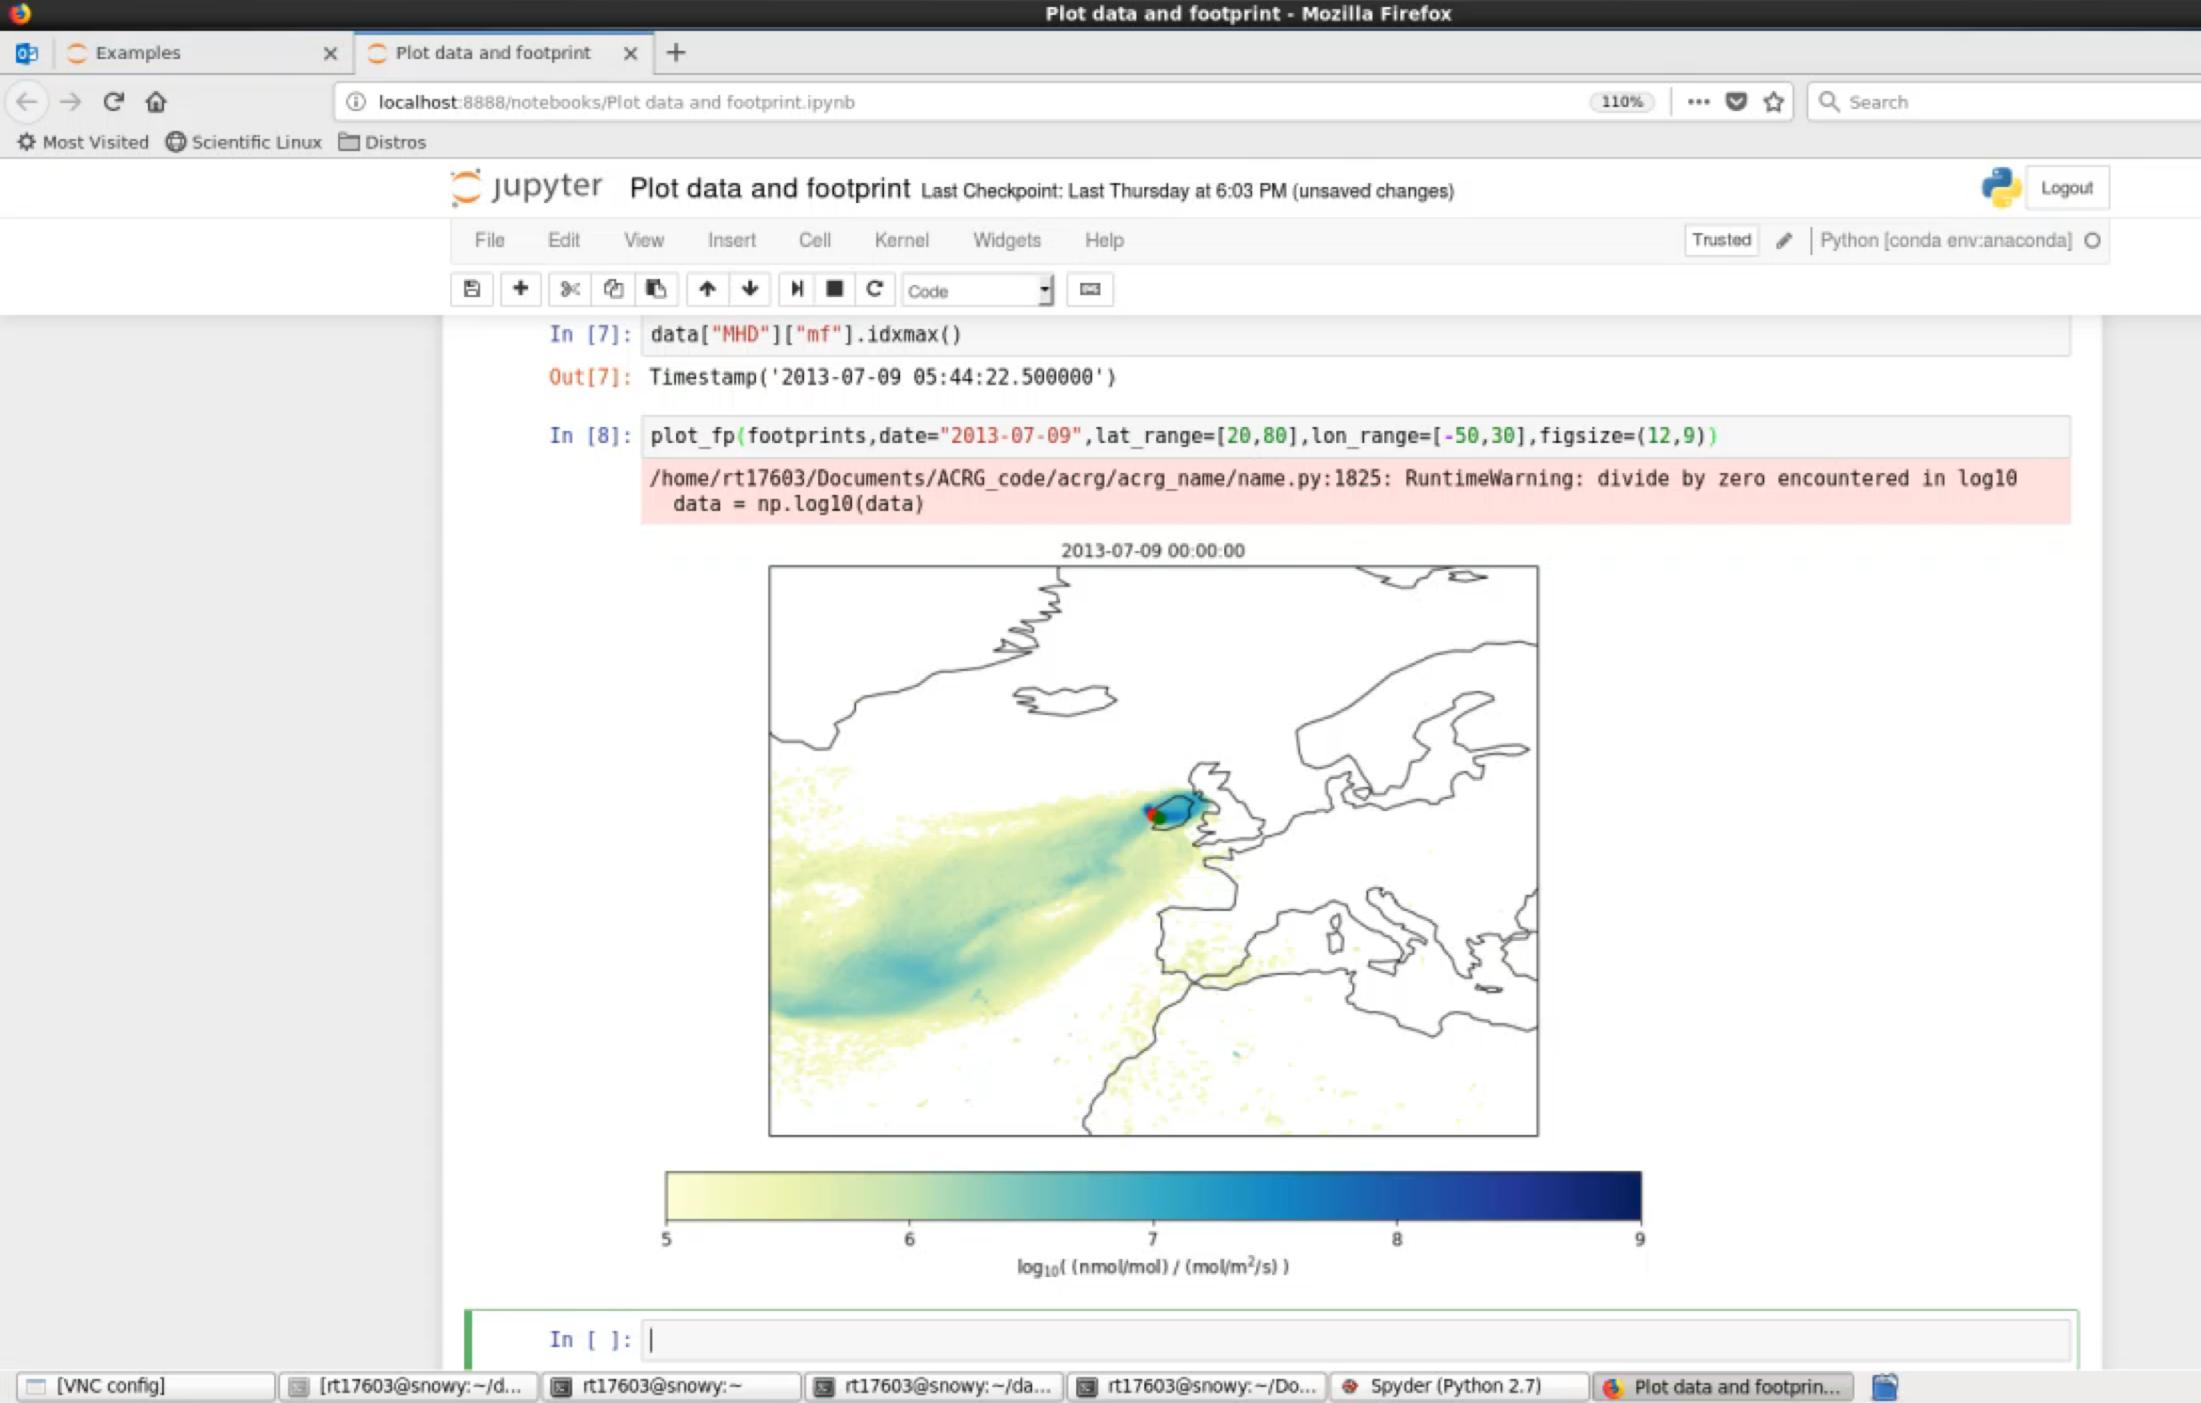


HUGS (Hub for UK Greenhouse Gas Science)

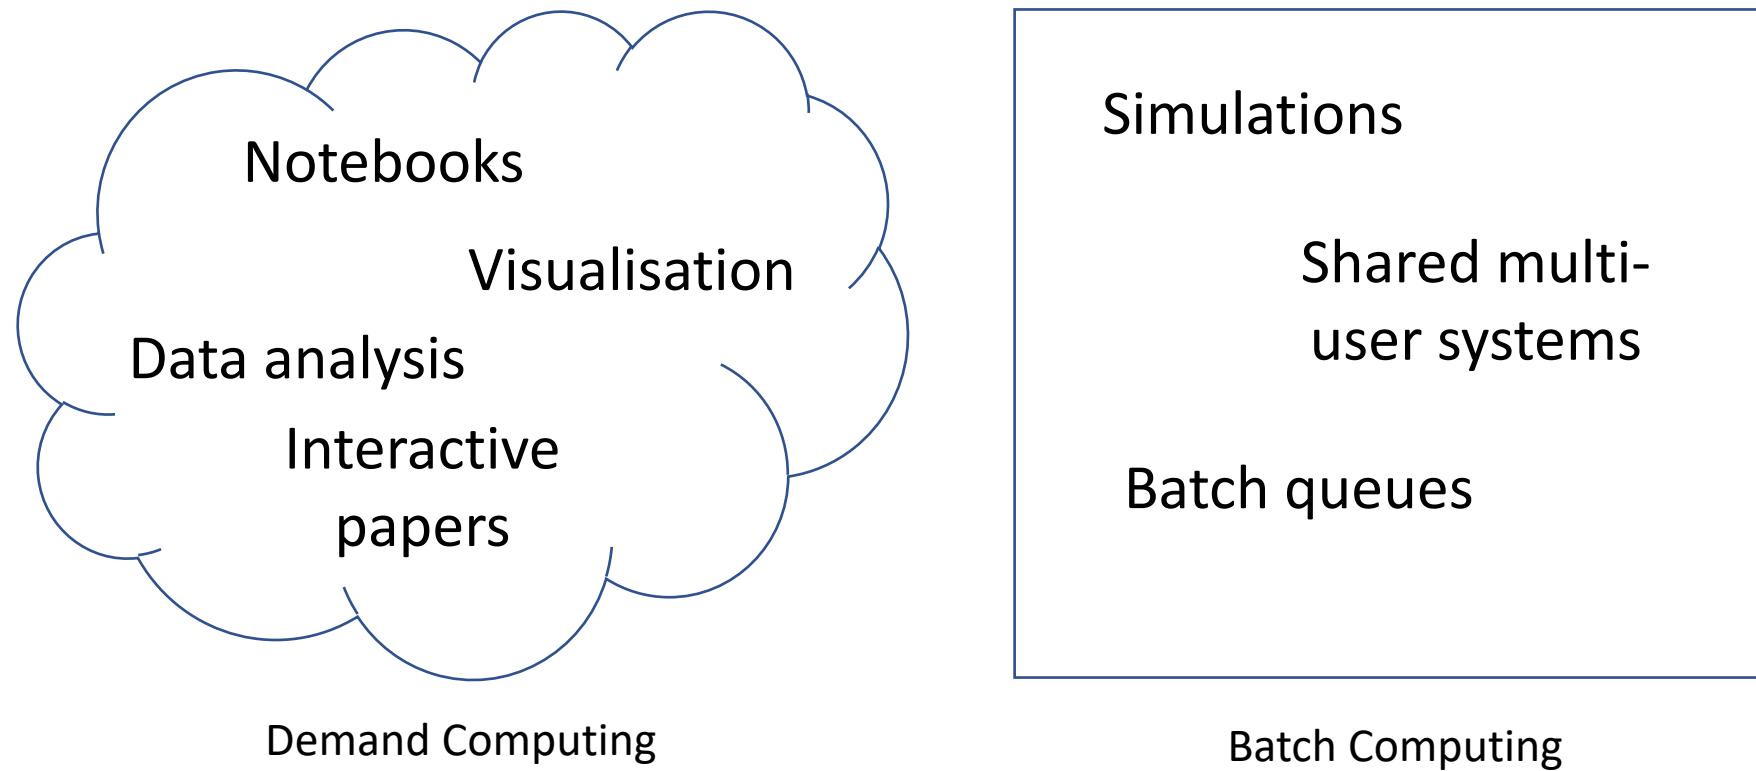


HUGS (Hub for UK Greenhouse Gas Science)

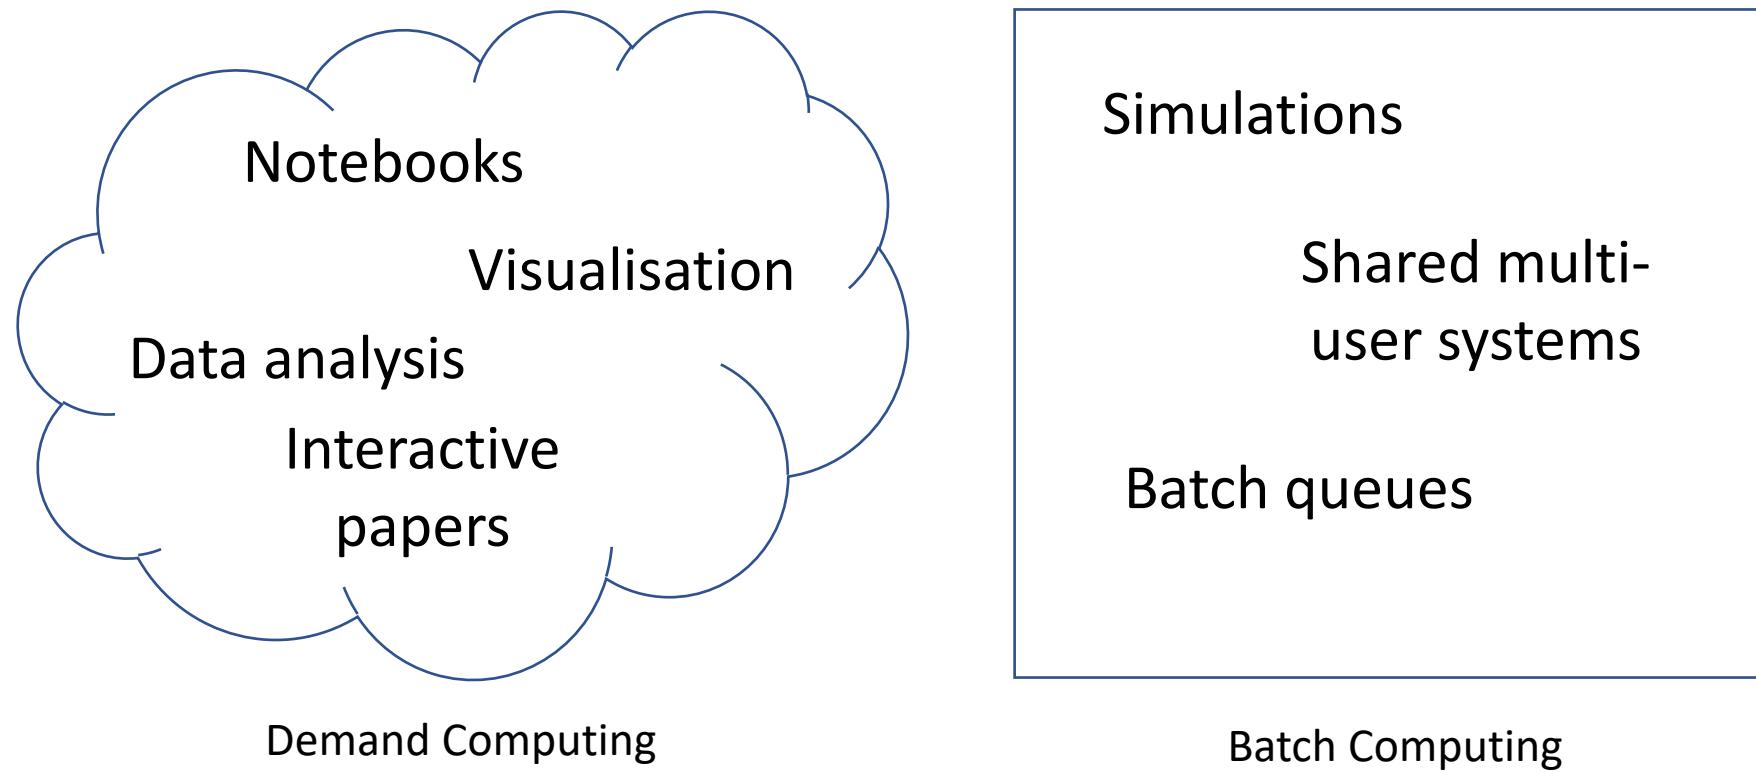




Demand versus Batch Computing



Demand versus Batch Computing



The Planetary Supercomputer

Functions == Processes



Object Store == Disk/Memory Storage

AAAI == User accounts and resource scheduler

Notebooks == Console/GUI

- Building a service that allows on-demand running of HPC workloads from within interactive Jupyter notebooks with a full user **Authentication, Access control and financial Accounting Infrastructure (AAAI)**
- Fn is an excellent function / serverless platform. Open source 😊
- Fully portable – works across clouds!
- Notebooks + Serverless + Object Store equals programming the planetary supercomputer
- Or, as my students call it, building the Netflix of Simulation

The Planetary Supercomputer

```
In [ ]: from Acquire.Client import User
In [ ]: user = User("chryswoods", identity_url="http://identity-gcp.acquire-aaai.com:8080/t/identity")
In [ ]: (url,qrcode) = user.request_login()
In [ ]: qrcode
In [ ]: from IPython.core.display import HTML
       print(url)
       HTML('<a href="%s">Login here</a>' % url)
In [ ]: user.wait_for_login()
In [ ]: user.is_logged_in()
In [ ]: user.logout()
In [ ]:
In [ ]:
```

```
In [ ]: from Acquire.Client import User
In [ ]: user = User("chryswoods", identity_url="http://identity-azure.acquire-aaai.com:8080/t/identity")
In [ ]: (url,qrcode) = user.request_login()
In [ ]: qrcode
In [ ]: from IPython.core.display import HTML
       print(url)
       HTML('<a href="%s">Login here</a>' % url)
In [ ]: user.wait_for_login()
In [ ]: user.is_logged_in()
In [ ]: user.logout()
In [ ]:
In [ ]:
```

```
In [ ]: from Acquire.Client import User
In [ ]: user = User("chryswoods", identity_url="http://identity.acquire-aaai.com:8080/t/identity")
In [ ]: (url,qrcode) = user.request_login()
In [ ]: qrcode
In [ ]: from IPython.core.display import HTML
       print(url)
       HTML('<a href="%s">Login here</a>' % url)
In [ ]: user.wait_for_login()
In [ ]: user.is_logged_in()
In [ ]: user.logout()
In [ ]:
In [ ]:
```



Fn running
on GCP in Japan



Notebooks running in Seattle



Fn running on Azure
in the Netherlands



OCI Object Store in Germany



Fn running on OCI
in Germany

Acknowledgements

My Research Software Engineering (RSE) Group

Andrew Williams, Chris Edsall, Lester Hedges, Matt Williams

BioSimSpace Research Team

Julien Michel, Antonia Mey, Adrian Mulholland, Charlie Laughton, Francesco Gervasio

HUGS Team

Matt Rigby, Simon O'Doherty, Rachel Tunnicliffe

EPSRC and NERC for funding (EP/N018591/1, EP/P022138/1, NE/S016155/1)

Oracle for providing a lot of compute time and extremely valuable discussions with cloud engineers and the Fn development team. Special thanks to Jenny Tsai-Smith, Phil Bates, Gerardo Viedma and Chad Arimura for all the help and useful technical discussions

Microsoft Azure (particularly Kenji Takeda and Mike Kiernan), and **Google** (particularly Ilias Katsardis) for compute time and engineer discussions to help develop and port to Azure and GCP

Amazon Web Services (particularly Francis Dauncey and Lucy Antysz) and **IBM** (particularly Michael Behrendt and Petrena Prince) for committing to providing compute time and engineer discussions to port to AWS and the IBM cloud

Identity, Access, Accounting, Storage, Compute

- Identity service supports secure login from any computer, notebook or cloud device. New key/cert generated for each login session and public key/cert uploaded to identity service. Accepted after out-of-band 2FA login. Used to sign-encrypt authorization tokens (json) for resources, actions and identifying users to Fn functions of other services
- Accounting service provides fully audited accounting and budget management. Users can create accounts with different daily spend limits, create and destroy accounts easily, create delegated accounts for members of their group etc. All transactions tracked and mandated via authorization tokens verified via trusted identity services
- Storage service provides fine-grained control (read/write) to objects or buckets to users via authorization tokens verified via trusted identity services. Bucket/object access is via a Fn-based API with encrypted OCI PARs used on the client and paid for using transaction credit notes and receipts processed via trusted accounting services
- Compute service (will) provision a dynamic cluster-in-the-cloud to allow any workload to run on any hardware, jobs described using json, data transferred via OCI PARs generated using the storage service, and authorized and paid for using transaction credit notes and receipts processed via trusted accounting services
- Access service (will) collect price lists from different trusted services and chooses which storage and compute services to delegate different work, and mediates payment for those services from users via authorization tokens validated by trusted identity services, and debit+credit notes from trusted accounting services, plus receipts on completion of work from trusted storage and compute services.

All built as distribute services built on top of Fn 😊